

#Automatic Port Operations

Author : Sundar Krishnamachari

Objectives

- Identify different classes of boats using a Covolutional Neural Network model
 - Identify different classes of boats using the MobileNet_V2 model using transfer learning
 - Compare the performanaces of both and draw inferences
-

Tasks performed on each model

1. Evaluate Model's performance by calculating accuracy and loss metrics
2. Make predictions on the test set
3. Print confusion matrix and display
4. Print classificaton report

##Common Functions

```
import os
import pathlib
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np
import seaborn as sns
import pandas as pd

from sklearn.metrics import confusion_matrix,
ConfusionMatrixDisplay,classification_report

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential

import warnings
warnings.filterwarnings('ignore')

train_accuracy = {}
test_accuracy = {}
train_loss = {}
test_loss = {}

def print_scores(model, val_ds, train_ds,name='Default'):
    """
    :param model:
    :param val_ds:
    :param train_ds:
```

```

: return:
"""
print(f'Displaying accuracy and loss for {name}')

score_test = model.evaluate(val_ds, verbose=0)
score_train = model.evaluate(train_ds, verbose=0)

print('Test loss: %.4f' % score_test[0])
print('Test accuracy: %.4f' % score_test[1])

print('Train loss: %.4f' % score_train[0])
print('train accuracy: %.4f' % score_train[1])

if not name == 'Default':
    test_loss[name] = score_test[0]
    test_accuracy[name] = score_test[1]
    train_loss[name] = score_train[0]
    train_accuracy[name] = score_train[1]

def _plot(history, param_1, param_2, ax):
    """
    :param history:
    :param param_1:
    :param param_2:
    :return:
    """

    ax.plot(history.history[param_1])
    ax.plot(history.history[param_2])
    ax.set_title(f'model {param_1}')
    ax.set_ylabel(param_1)
    ax.set_xlabel('epoch')
    ax.legend(['train', 'test'], loc='best')

def plot_accuracy_loss_graphs(history, title=None, show=True):
    """
    :param history:
    :param title:
    :param show:
    :return:
    """

    print()
    fig = plt.figure(figsize=(10, 10))

    # accuracy chart
    ax = plt.subplot(1, 2, 1)

```

```

    _plot(history, 'accuracy', 'val_accuracy', ax)
    # loss chart
    ax = plt.subplot(1, 2, 2)
    _plot(history, 'loss', 'val_loss', ax)
    if show:
        plt.tight_layout()
        plt.show()
    else:
        plt.savefig(f'{title}.png')
        plt.close()

def get_predicted_labels(model, class_names):
    """
    :param model:
    :param class_names:
    :return: dict of image paths and image labels
    """
    images = os.listdir(test_dir)
    image_label_map = {}

    for i, image in enumerate(images):
        path = test_dir + '/' + image
        img = tf.keras.utils.load_img(path, target_size=(height,
width))
        img_array = tf.keras.utils.img_to_array(img)
        img_array = tf.expand_dims(img_array, 0) # Create a batch
        predictions = model.predict(img_array, verbose=False)
        score = tf.nn.softmax(predictions[0])
        max_score = np.max(score)
        image_label_map[path] = class_names[np.argmax(score)]
    return image_label_map

def display_test_images(image_map, model_name):
    """
    :param image_map:
    :param model_name:
    :return:
    """
    print(f'Predictions for {model_name}\n')
    plt.figure(figsize=(25, 25))
    for i, image_path in enumerate(image_map.keys()):
        ax = plt.subplot(20, 15, i + 1)
        img = mpimg.imread(image_path)
        ax.imshow(img)
        ax.axis('off')
        label = image_map.get(image_path)

```

```

        # ax.text(0.5, .8, label, fontsize=10, fontweight='bold',
        color='green', ha='center', va='center', transform=ax.transAxes)
        ax.set_title(label, fontsize=10)
        plt.show()

def display_cm(true_labels, predicted_labels):
    """
    :param true_labels:
    :param predicted_labels:
    :return:
    """
    cm = confusion_matrix(true_labels, predicted_labels)
    print('Confusion Matrix')
    print(cm)
    print()

    print('Confusion Matrix Displayed\n')
    fig = plt.figure(figsize=(10, 10))
    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
    display_labels=class_names)
    disp.plot()
    plt.title('Confusion Matrix')
    plt.xticks(rotation=90)
    plt.show()

def get_metrics(model, validation_ds, val_split=0.2):
    num_images = int(total_train_images * val_split)
    num_batches = int(num_images / 32)
    residuals = int(num_images % 32)
    true_labels = []
    predicted_labels = []
    for images, labels in validation_ds.take(num_batches):
        true_labels.extend(labels)
        for i in range(batch_size): # prediction for each image in
the batch
            img_array = tf.keras.utils.img_to_array(images[i])
            img_array = tf.expand_dims(img_array, 0) # Create a batch
            predictions = model.predict(img_array, verbose=False)
            score = tf.nn.softmax(predictions[0])
            predicted_labels.append(np.argmax(score))
    return true_labels, predicted_labels

def create_transfer_model(transfer_model, dropout=0.1,
trainable=False, data_augmentation=None):
    transfer_model.trainable = trainable
    model_t = Sequential()
    if not data_augmentation is None:

```

```

        for aug in data_augmentation:
            print(f'added {aug} to model')
            model_t.add(aug)
    if data_augmentation is None:
        # data augmentation will add an input shape in the first layer
        model_t.add(layers.Rescaling(1. / 255, input_shape=(height,
width, 3)))
    else:
        model_t.add(layers.Rescaling(1. / 255))
    model_t.add(transfer_model)
    model_t.add(layers.GlobalAveragePooling2D())
    model_t.add(layers.Dropout(dropout))

    # FCN
    model_t.add(layers.Flatten())
    model_t.add(layers.Dense(256, activation='relu'))
    model_t.add(layers.BatchNormalization())
    model_t.add(layers.Dropout(dropout))

    model_t.add(layers.Dense(128, activation='relu'))
    model_t.add(layers.BatchNormalization())
    model_t.add(layers.Dropout(dropout))

    model_t.add(layers.Dense(len(class_names), activation='softmax'))
    model_t.summary()
    return model_t

```

##CNN Model

##Load Datasets

```

# from google.colab import drive
# drive.mount('/content/drive')

```

Mounted at /content/drive

##Extract data

```

# from zipfile import ZipFile

# zip_ref = ZipFile('/content/drive/MyDrive/DL/BOATS.zip', 'r')
# zip_ref.extractall('/content/drive/MyDrive/DL')
# zip_ref.close()

```

##Define constants , test and train directories

```

# lets define some constants
train_dir_kaggle = '/kaggle/input/boats-ds/TRAIN_BOATS'
test_dir_kaggle = '/kaggle/input/boats-ds/TEST_BOATS'

train_dir_colab = '/content/drive/MyDrive/DL/TRAIN_BOATS/'
test_dir_colab = '/content/drive/MyDrive/DL/TEST_BOATS/'

train_dir = train_dir_colab
test_dir = test_dir_colab

if os.path.exists(train_dir_kaggle):
    train_dir = train_dir_kaggle
    test_dir = test_dir_kaggle
    print('Using Kaggle data')
    print(f'train_dir = {train_dir}')
    print(f'test_dir = {test_dir}')
else:
    print('Using Colab data')
    print(f'train_dir = {train_dir}')
    print(f'test_dir = {test_dir}')

height = 224
width = 224
batch_size = 32
total_train_images = 0

total_test_images = len(os.listdir(test_dir))

for sub_dir in os.listdir(train_dir):
    total_train_images += len(os.listdir(train_dir+'/'+sub_dir))

print('total train images = ', total_train_images)
print('total test images = ', total_test_images)

Using Colab data
train_dir = /content/drive/MyDrive/DL/TRAIN_BOATS/
test_dir = /content/drive/MyDrive/DL/TEST_BOATS/
total train images = 1162
total test images = 300

```

##Estimate number of images per class and plot

```

data_dir = pathlib.Path(train_dir)

def get_subdirectories(directory):
    with os.scandir(directory) as entries:

```

```

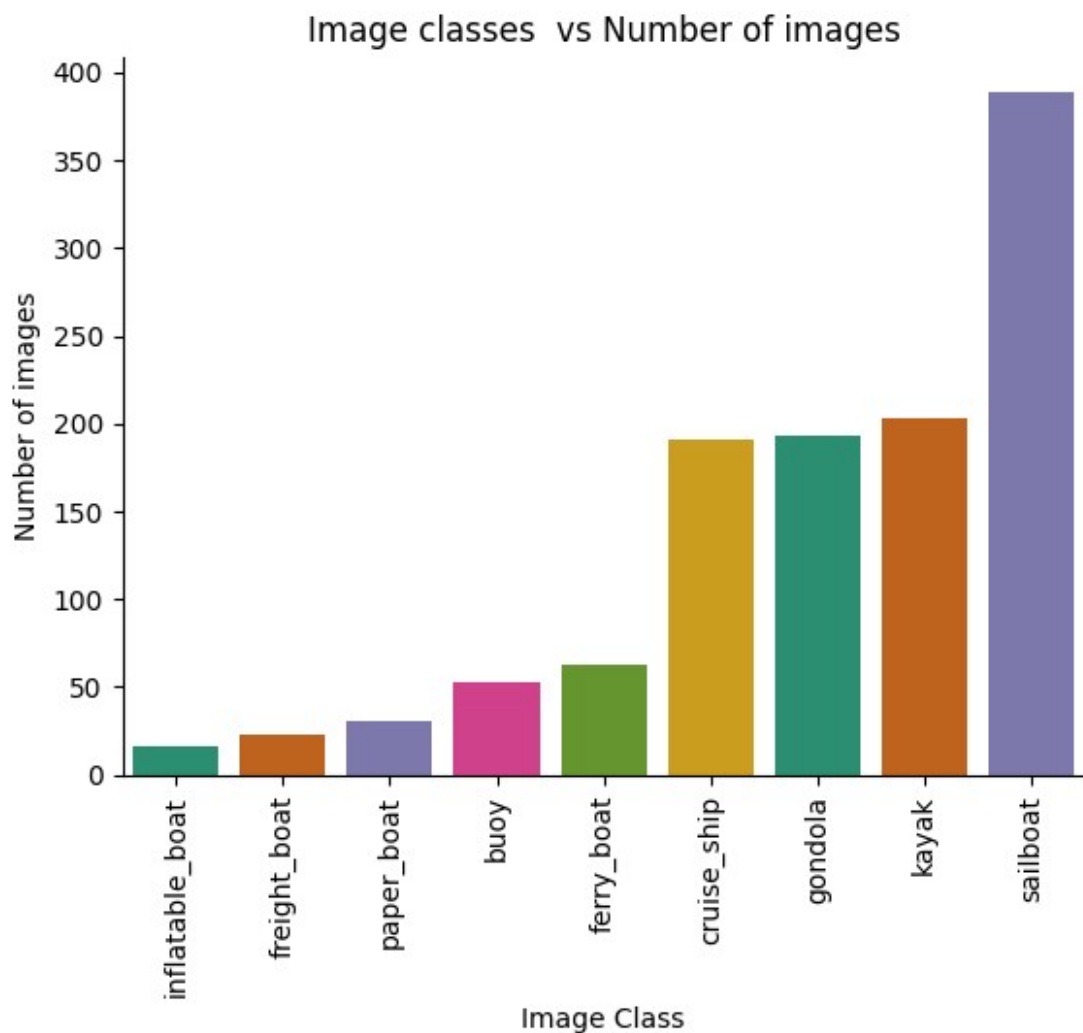
        return [entry.name for entry in entries if entry.is_dir()]
sub_dirs = get_subdirectories(data_dir)
images = {}
image_class = []
num_images = []
class_dict = {}

for d in sub_dirs:
    image_class.append(d)
    num_images.append(len(os.listdir(train_dir+'/'+d)))

df =
pd.DataFrame({'Image_Class':image_class,'Num_Images':num_images}).sort_
_values(by='Num_Images',ascending=True).reset_index()

indices = sorted(df['index'].values.tolist())
class_names = df['Image_Class'].values
for i in indices:
    class_dict[i] = class_names[i]
sns.barplot(data=df,x='Image_Class',y='Num_Images',palette=sns.mpl_pal
ette('Dark2'))
plt.xlabel('Image Class')
plt.ylabel('Number of images')
plt.title(f'Image classes vs Number of images' )
plt.xticks(rotation = 90)
plt.gca().spines[['top', 'right']].set_visible(False)
plt.show()
df

```



```
{
  "summary": {
    "name": "df",
    "rows": 9,
    "fields": [
      {
        "column": "index",
        "properties": {
          "dtype": "number",
          "std": 2,
          "min": 0,
          "max": 8,
          "num_unique_values": 9,
          "samples": [
            6, 3, 1
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "Image_Class",
        "properties": {
          "dtype": "string",
          "num_unique_values": 9,
          "samples": [
            "kayak", "freight_boat", "cruise_ship"
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "Num_Images",
        "properties": {
          "dtype": "number",
          "std": 124,
          "min": 16,
          "max": 389,
          "num_unique_values": 9,
          "samples": [
            203, 23, 191
          ],
          "semantic_type": "",
          "description": ""
        }
      }
    ]
  },
  "type": "dataframe",
  "variable_name": "df"
}
```


There is a predominance of the sailboat class in the training data. The bottom three image classes have less than 50 images per class. Ideally we can drop them. Having them in the training/testing datasets can skew the predictions. But, I am keeping them nonetheless.

##Define datasets for training and validating the model

```
# Not using ImageDataGenerator as it is deprecated and it doesn't allow for  
# caching and prefetching. With image data, caching and prefetching can  
determine  
# if training will be possible with limited gpu resources  
  
# https://www.tensorflow.org/guide/data  
  
import tensorflow as tf  
  
def get_ds(data_dir, validation_split=0.2, subset='training', seed =  
43, image_size=(255, 255), batch_size=32):  
    train_ds = tf.keras.utils.image_dataset_from_directory(  
        data_dir,  
        validation_split=validation_split,  
        subset=subset,  
        seed=seed,  
        image_size=image_size,  
        batch_size=batch_size)  
    train_ds.class_names  
    return train_ds  
  
train_ds =  
get_ds(data_dir, image_size=(height, width), batch_size=batch_size)  
val_ds = get_ds(data_dir, validation_split=0.2, subset='validation', seed  
= 43, image_size=(height, width), batch_size=batch_size)  
class_names = train_ds.class_names  
print(class_names)  
  
Found 1162 files belonging to 9 classes.  
Using 930 files for training.  
Found 1162 files belonging to 9 classes.  
Using 232 files for validation.  
['buoy', 'cruise_ship', 'ferry_boat', 'freight_boat', 'gondola',  
'inflatable_boat', 'kayak', 'paper_boat', 'sailboat']
```

##Plot some images from the training dataset

```
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(10, 10))  
  
for images, labels in train_ds.take(1):  
    for i in range(9):
```

```
ax = plt.subplot(3, 3, i + 1)
plt.imshow(images[i].numpy().astype("uint8"))
plt.title(class_names[labels[i]])
plt.axis("off")
```

cruise_ship



kayak



kayak



kayak



kayak




cruise_ship




gondola



gondola



sailboat



##Set up the datasets for caching and prefetching

```
# Using `cache`, `prfetch` and `AUTOTUNE` for efficient handling of
input data.
# https://www.tensorflow.org/guide/data\_performance
AUTOTUNE = tf.data.AUTOTUNE
```

```

train_ds =
train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

##CNN model exactly as specified in the problem 2 Conv2d layers with Max 2d pooling

1 2D Global Average Pooling

2 FFN with 2 dense layers of 128 neurons each

1 output layer with 9 neurons and softmax activation

```

model = Sequential([
    # Image Scaling layer
    layers.Rescaling(1./255, input_shape=(height, width, 3)),
    # First convolution layer.convolution with 32 filters with a 3x3
    kernel matrix,no padding
    layers.Conv2D(32, (3,3), activation='relu'),
    layers.MaxPooling2D(),
    #layers.Dropout(0.1),
    # Second Convolution layer
    layers.Conv2D(32, (3,3), activation='relu'),
    layers.MaxPooling2D(),
    #Global Average Pooling
    layers.GlobalAveragePooling2D(),
    # FFN
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(128, activation='relu'),
    layers.Dense(len(class_names),activation='softmax')
])

```

##Compile the model

```

model.compile(optimizer='adam',
#https://www.tensorflow.org/api_docs/python/tf/keras/losses/CategoricalCrossentropy
#It states:

```

```

        # "We expect labels to be provided in a one_hot
representation.
        # If you want to provide labels as integers, please use
SparseCategoricalCrossentropy loss."
        loss = 'sparse_categorical_crossentropy',
        # keras has removed support for batch level precision
and recall metrics.Problem statement is obsolete
        metrics=['accuracy']
    )
model.summary()
Model: "sequential"

```

Layer (type) Param #	Output Shape
rescaling (Rescaling) 0	(None, 224, 224, 3)
conv2d (Conv2D) 896	(None, 222, 222, 32)
max_pooling2d (MaxPooling2D) 0	(None, 111, 111, 32)
conv2d_1 (Conv2D) 9,248	(None, 109, 109, 32)
max_pooling2d_1 (MaxPooling2D) 0	(None, 54, 54, 32)
global_average_pooling2d (GlobalAveragePooling2D) 0	(None, 32)
flatten (Flatten) 0	(None, 32)
dense (Dense)	(None, 128)

4,224		
dense_1 (Dense)	(None, 128)	
16,512		
dense_2 (Dense)	(None, 9)	
1,161		

Total params: 32,041 (125.16 KB)

Trainable params: 32,041 (125.16 KB)

```
Non-trainable params: 0 (0.00 B)
```

Train the model with Early Stopping and Checkpoint

```
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

#Optimizing for maximum accuracy instead of minimum loss.( min loss does not always mean max accuracy)
early_stop=EarlyStopping(monitor='val_accuracy',mode='max',verbose=1,patience=10,min_delta=0.01)
checkpoint=ModelCheckpoint('best_model.keras',monitor='val_accuracy',mode='max',verbose=1,save_best_only=True)

epochs=100
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    batch_size = batch_size,
    # not specified. But i thought this was a good idea
    callbacks = [early_stop,checkpoint]
)

Epoch 1/100
29/30 ————— 0s 79ms/step - accuracy: 0.2852 - loss: 2.0488
Epoch 1: val_accuracy improved from -inf to 0.30172, saving model to best_model.keras
30/30 ————— 185s 1s/step - accuracy: 0.2881 - loss: 2.0421 - val_accuracy: 0.3017 - val_loss: 1.8125
Epoch 2/100
30/30 ————— 0s 27ms/step - accuracy: 0.3327 - loss: 1.8598
Epoch 2: val accuracy did not improve from 0.30172
```

```
30/30 _____ 1s 30ms/step - accuracy: 0.3330 - loss:
1.8580 - val_accuracy: 0.3017 - val_loss: 1.7931
Epoch 3/100
30/30 _____ 0s 26ms/step - accuracy: 0.3450 - loss:
1.7445
Epoch 3: val_accuracy did not improve from 0.30172
30/30 _____ 1s 28ms/step - accuracy: 0.3449 - loss:
1.7460 - val_accuracy: 0.3017 - val_loss: 1.7872
Epoch 4/100
30/30 _____ 0s 26ms/step - accuracy: 0.3387 - loss:
1.7577
Epoch 4: val_accuracy did not improve from 0.30172
30/30 _____ 1s 31ms/step - accuracy: 0.3390 - loss:
1.7580 - val_accuracy: 0.3017 - val_loss: 1.7648
Epoch 5/100
30/30 _____ 0s 25ms/step - accuracy: 0.3479 - loss:
1.8035
Epoch 5: val_accuracy improved from 0.30172 to 0.31034, saving model
to best_model.keras
30/30 _____ 1s 28ms/step - accuracy: 0.3485 - loss:
1.8015 - val_accuracy: 0.3103 - val_loss: 1.7453
Epoch 6/100
30/30 _____ 0s 25ms/step - accuracy: 0.3658 - loss:
1.6941
Epoch 6: val_accuracy improved from 0.31034 to 0.31897, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.3661 - loss:
1.6944 - val_accuracy: 0.3190 - val_loss: 1.7383
Epoch 7/100
30/30 _____ 0s 25ms/step - accuracy: 0.4214 - loss:
1.6582
Epoch 7: val_accuracy did not improve from 0.31897
30/30 _____ 1s 27ms/step - accuracy: 0.4206 - loss:
1.6593 - val_accuracy: 0.3060 - val_loss: 1.7430
Epoch 8/100
29/30 _____ 0s 25ms/step - accuracy: 0.3546 - loss:
1.6897
Epoch 8: val_accuracy improved from 0.31897 to 0.32759, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.3574 - loss:
1.6887 - val_accuracy: 0.3276 - val_loss: 1.7422
Epoch 9/100
30/30 _____ 0s 25ms/step - accuracy: 0.4039 - loss:
1.6579
Epoch 9: val_accuracy did not improve from 0.32759
30/30 _____ 1s 27ms/step - accuracy: 0.4041 - loss:
1.6581 - val_accuracy: 0.3103 - val_loss: 1.7343
Epoch 10/100
30/30 _____ 0s 25ms/step - accuracy: 0.3987 - loss:
```

```
1.6922
Epoch 10: val_accuracy did not improve from 0.32759
30/30 ━━━━━━━━━━━ 1s 27ms/step - accuracy: 0.3988 - loss:
1.6909 - val_accuracy: 0.3276 - val_loss: 1.7060
Epoch 11/100
30/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.4005 - loss:
1.6570
Epoch 11: val_accuracy did not improve from 0.32759
30/30 ━━━━━━━━━━━ 1s 27ms/step - accuracy: 0.4006 - loss:
1.6571 - val_accuracy: 0.3190 - val_loss: 1.8027
Epoch 12/100
29/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.4045 - loss:
1.6640
Epoch 12: val_accuracy did not improve from 0.32759
30/30 ━━━━━━━━━━━ 1s 27ms/step - accuracy: 0.4052 - loss:
1.6630 - val_accuracy: 0.3276 - val_loss: 1.7219
Epoch 13/100
30/30 ━━━━━━━━━━━ 0s 26ms/step - accuracy: 0.3792 - loss:
1.6355
Epoch 13: val_accuracy did not improve from 0.32759
30/30 ━━━━━━━━━━━ 1s 28ms/step - accuracy: 0.3802 - loss:
1.6355 - val_accuracy: 0.3190 - val_loss: 1.7777
Epoch 14/100
28/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.4101 - loss:
1.6534
Epoch 14: val_accuracy improved from 0.32759 to 0.33190, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 29ms/step - accuracy: 0.4108 - loss:
1.6535 - val_accuracy: 0.3319 - val_loss: 1.7044
Epoch 15/100
29/30 ━━━━━━━━━━━ 0s 26ms/step - accuracy: 0.4504 - loss:
1.5937
Epoch 15: val_accuracy improved from 0.33190 to 0.34052, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 33ms/step - accuracy: 0.4486 - loss:
1.5961 - val_accuracy: 0.3405 - val_loss: 1.6975
Epoch 16/100
30/30 ━━━━━━━━━━━ 0s 26ms/step - accuracy: 0.4283 - loss:
1.6227
Epoch 16: val_accuracy did not improve from 0.34052
30/30 ━━━━━━━━━━━ 1s 28ms/step - accuracy: 0.4281 - loss:
1.6224 - val_accuracy: 0.3319 - val_loss: 1.6979
Epoch 17/100
30/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.4119 - loss:
1.6808
Epoch 17: val_accuracy did not improve from 0.34052
30/30 ━━━━━━━━━━━ 1s 28ms/step - accuracy: 0.4122 - loss:
1.6790 - val_accuracy: 0.3233 - val_loss: 1.7019
Epoch 18/100
```

```
29/30 _____ 0s 25ms/step - accuracy: 0.4370 - loss:
1.6216
Epoch 18: val_accuracy did not improve from 0.34052
30/30 _____ 1s 27ms/step - accuracy: 0.4366 - loss:
1.6203 - val_accuracy: 0.3319 - val_loss: 1.7042
Epoch 19/100
30/30 _____ 0s 25ms/step - accuracy: 0.4183 - loss:
1.6161
Epoch 19: val_accuracy improved from 0.34052 to 0.34914, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.4184 - loss:
1.6155 - val_accuracy: 0.3491 - val_loss: 1.6773
Epoch 20/100
30/30 _____ 0s 25ms/step - accuracy: 0.4340 - loss:
1.5819
Epoch 20: val_accuracy did not improve from 0.34914
30/30 _____ 1s 27ms/step - accuracy: 0.4336 - loss:
1.5824 - val_accuracy: 0.3319 - val_loss: 1.6749
Epoch 21/100
30/30 _____ 0s 25ms/step - accuracy: 0.4088 - loss:
1.6157
Epoch 21: val_accuracy improved from 0.34914 to 0.36207, saving model
to best_model.keras
30/30 _____ 1s 28ms/step - accuracy: 0.4096 - loss:
1.6151 - val_accuracy: 0.3621 - val_loss: 1.6463
Epoch 22/100
30/30 _____ 0s 25ms/step - accuracy: 0.4592 - loss:
1.5375
Epoch 22: val_accuracy did not improve from 0.36207
30/30 _____ 1s 27ms/step - accuracy: 0.4584 - loss:
1.5387 - val_accuracy: 0.3405 - val_loss: 1.6630
Epoch 23/100
30/30 _____ 0s 25ms/step - accuracy: 0.4134 - loss:
1.5790
Epoch 23: val_accuracy did not improve from 0.36207
30/30 _____ 1s 27ms/step - accuracy: 0.4141 - loss:
1.5786 - val_accuracy: 0.3578 - val_loss: 1.6266
Epoch 24/100
30/30 _____ 0s 25ms/step - accuracy: 0.4524 - loss:
1.5239
Epoch 24: val_accuracy did not improve from 0.36207
30/30 _____ 1s 27ms/step - accuracy: 0.4519 - loss:
1.5251 - val_accuracy: 0.3578 - val_loss: 1.6500
Epoch 25/100
30/30 _____ 0s 25ms/step - accuracy: 0.4447 - loss:
1.5227
Epoch 25: val_accuracy did not improve from 0.36207
30/30 _____ 1s 27ms/step - accuracy: 0.4447 - loss:
1.5235 - val_accuracy: 0.3362 - val_loss: 1.6190
```



```
Epoch 26/100
30/30 _____ 0s 25ms/step - accuracy: 0.4715 - loss:
1.4954
Epoch 26: val_accuracy improved from 0.36207 to 0.37500, saving model
to best_model.keras
30/30 _____ 1s 28ms/step - accuracy: 0.4711 - loss:
1.4968 - val_accuracy: 0.3750 - val_loss: 1.5957
Epoch 27/100
30/30 _____ 0s 25ms/step - accuracy: 0.4502 - loss:
1.5242
Epoch 27: val_accuracy did not improve from 0.37500
30/30 _____ 1s 27ms/step - accuracy: 0.4502 - loss:
1.5248 - val_accuracy: 0.3578 - val_loss: 1.6842
Epoch 28/100
30/30 _____ 0s 25ms/step - accuracy: 0.4747 - loss:
1.5090
Epoch 28: val_accuracy improved from 0.37500 to 0.38362, saving model
to best_model.keras
30/30 _____ 1s 28ms/step - accuracy: 0.4740 - loss:
1.5099 - val_accuracy: 0.3836 - val_loss: 1.5858
Epoch 29/100
29/30 _____ 0s 25ms/step - accuracy: 0.4512 - loss:
1.4877
Epoch 29: val_accuracy did not improve from 0.38362
30/30 _____ 1s 27ms/step - accuracy: 0.4513 - loss:
1.4898 - val_accuracy: 0.3534 - val_loss: 1.5949
Epoch 30/100
30/30 _____ 0s 25ms/step - accuracy: 0.4364 - loss:
1.5884
Epoch 30: val_accuracy did not improve from 0.38362
30/30 _____ 1s 28ms/step - accuracy: 0.4374 - loss:
1.5866 - val_accuracy: 0.3750 - val_loss: 1.5709
Epoch 31/100
30/30 _____ 0s 25ms/step - accuracy: 0.4769 - loss:
1.4901
Epoch 31: val_accuracy improved from 0.38362 to 0.41379, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.4765 - loss:
1.4904 - val_accuracy: 0.4138 - val_loss: 1.6129
Epoch 32/100
30/30 _____ 0s 25ms/step - accuracy: 0.4549 - loss:
1.5267
Epoch 32: val_accuracy did not improve from 0.41379
30/30 _____ 1s 27ms/step - accuracy: 0.4559 - loss:
1.5256 - val_accuracy: 0.3750 - val_loss: 1.5965
Epoch 33/100
30/30 _____ 0s 25ms/step - accuracy: 0.4969 - loss:
1.4395
Epoch 33: val_accuracy did not improve from 0.41379
```

```
30/30 _____ 1s 27ms/step - accuracy: 0.4967 - loss:
1.4407 - val_accuracy: 0.4138 - val_loss: 1.5971
Epoch 34/100
30/30 _____ 0s 25ms/step - accuracy: 0.5139 - loss:
1.4200
Epoch 34: val_accuracy did not improve from 0.41379
30/30 _____ 1s 27ms/step - accuracy: 0.5133 - loss:
1.4215 - val_accuracy: 0.3707 - val_loss: 1.5685
Epoch 35/100
30/30 _____ 0s 25ms/step - accuracy: 0.4459 - loss:
1.5455
Epoch 35: val_accuracy did not improve from 0.41379
30/30 _____ 1s 27ms/step - accuracy: 0.4466 - loss:
1.5435 - val_accuracy: 0.3966 - val_loss: 1.5519
Epoch 36/100
30/30 _____ 0s 25ms/step - accuracy: 0.4704 - loss:
1.4832
Epoch 36: val_accuracy improved from 0.41379 to 0.42241, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.4708 - loss:
1.4827 - val_accuracy: 0.4224 - val_loss: 1.5245
Epoch 37/100
30/30 _____ 0s 25ms/step - accuracy: 0.5102 - loss:
1.4282
Epoch 37: val_accuracy improved from 0.42241 to 0.45259, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.5096 - loss:
1.4295 - val_accuracy: 0.4526 - val_loss: 1.5232
Epoch 38/100
30/30 _____ 0s 25ms/step - accuracy: 0.4749 - loss:
1.4864
Epoch 38: val_accuracy did not improve from 0.45259
30/30 _____ 1s 27ms/step - accuracy: 0.4751 - loss:
1.4863 - val_accuracy: 0.4267 - val_loss: 1.5239
Epoch 39/100
30/30 _____ 0s 25ms/step - accuracy: 0.4659 - loss:
1.4905
Epoch 39: val_accuracy improved from 0.45259 to 0.45690, saving model
to best_model.keras
30/30 _____ 1s 29ms/step - accuracy: 0.4669 - loss:
1.4891 - val_accuracy: 0.4569 - val_loss: 1.5199
Epoch 40/100
30/30 _____ 0s 25ms/step - accuracy: 0.5146 - loss:
1.4142
Epoch 40: val_accuracy did not improve from 0.45690
30/30 _____ 1s 27ms/step - accuracy: 0.5141 - loss:
1.4149 - val_accuracy: 0.4095 - val_loss: 1.5277
Epoch 41/100
30/30 _____ 0s 25ms/step - accuracy: 0.5006 - loss:
```

```
1.4690
Epoch 41: val_accuracy did not improve from 0.45690
30/30 ━━━━━━━━━━━ 1s 28ms/step - accuracy: 0.5006 - loss:
1.4683 - val_accuracy: 0.4138 - val_loss: 1.6197
Epoch 42/100
30/30 ━━━━━━━━━━━ 0s 26ms/step - accuracy: 0.5218 - loss:
1.4158
Epoch 42: val_accuracy improved from 0.45690 to 0.46121, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 30ms/step - accuracy: 0.5211 - loss:
1.4166 - val_accuracy: 0.4612 - val_loss: 1.5268
Epoch 43/100
30/30 ━━━━━━━━━━━ 0s 26ms/step - accuracy: 0.5019 - loss:
1.4511
Epoch 43: val_accuracy did not improve from 0.46121
30/30 ━━━━━━━━━━━ 1s 31ms/step - accuracy: 0.5023 - loss:
1.4502 - val_accuracy: 0.4224 - val_loss: 1.5743
Epoch 44/100
28/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.4883 - loss:
1.4382
Epoch 44: val_accuracy did not improve from 0.46121
30/30 ━━━━━━━━━━━ 1s 27ms/step - accuracy: 0.4898 - loss:
1.4359 - val_accuracy: 0.4440 - val_loss: 1.5133
Epoch 45/100
30/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.5217 - loss:
1.4347
Epoch 45: val_accuracy did not improve from 0.46121
30/30 ━━━━━━━━━━━ 1s 27ms/step - accuracy: 0.5215 - loss:
1.4346 - val_accuracy: 0.4052 - val_loss: 1.6003
Epoch 46/100
30/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.5069 - loss:
1.4543
Epoch 46: val_accuracy improved from 0.46121 to 0.46983, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 29ms/step - accuracy: 0.5068 - loss:
1.4543 - val_accuracy: 0.4698 - val_loss: 1.5034
Epoch 47/100
30/30 ━━━━━━━━━━━ 0s 26ms/step - accuracy: 0.5216 - loss:
1.4135
Epoch 47: val_accuracy improved from 0.46983 to 0.48276, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 30ms/step - accuracy: 0.5210 - loss:
1.4139 - val_accuracy: 0.4828 - val_loss: 1.4841
Epoch 48/100
29/30 ━━━━━━━━━━━ 0s 25ms/step - accuracy: 0.5428 - loss:
1.3631
Epoch 48: val_accuracy did not improve from 0.48276
30/30 ━━━━━━━━━━━ 1s 27ms/step - accuracy: 0.5413 - loss:
1.3651 - val_accuracy: 0.4784 - val_loss: 1.4935
```

```
Epoch 49/100
30/30 _____ 0s 25ms/step - accuracy: 0.5158 - loss:
1.3895
Epoch 49: val_accuracy did not improve from 0.48276
30/30 _____ 1s 28ms/step - accuracy: 0.5161 - loss:
1.3895 - val_accuracy: 0.4569 - val_loss: 1.5253
Epoch 50/100
30/30 _____ 0s 25ms/step - accuracy: 0.5049 - loss:
1.3991
Epoch 50: val_accuracy did not improve from 0.48276
30/30 _____ 1s 27ms/step - accuracy: 0.5050 - loss:
1.3996 - val_accuracy: 0.4784 - val_loss: 1.4883
Epoch 51/100
30/30 _____ 0s 25ms/step - accuracy: 0.5082 - loss:
1.4229
Epoch 51: val_accuracy did not improve from 0.48276
30/30 _____ 1s 28ms/step - accuracy: 0.5085 - loss:
1.4217 - val_accuracy: 0.4741 - val_loss: 1.4719
Epoch 52/100
29/30 _____ 0s 26ms/step - accuracy: 0.5348 - loss:
1.3502
Epoch 52: val_accuracy did not improve from 0.48276
30/30 _____ 1s 28ms/step - accuracy: 0.5328 - loss:
1.3539 - val_accuracy: 0.4267 - val_loss: 1.5252
Epoch 53/100
30/30 _____ 0s 25ms/step - accuracy: 0.5221 - loss:
1.4030
Epoch 53: val_accuracy did not improve from 0.48276
30/30 _____ 1s 27ms/step - accuracy: 0.5218 - loss:
1.4028 - val_accuracy: 0.4483 - val_loss: 1.5108
Epoch 54/100
30/30 _____ 0s 25ms/step - accuracy: 0.5228 - loss:
1.4032
Epoch 54: val_accuracy did not improve from 0.48276
30/30 _____ 1s 31ms/step - accuracy: 0.5224 - loss:
1.4035 - val_accuracy: 0.4483 - val_loss: 1.5235
Epoch 55/100
30/30 _____ 0s 26ms/step - accuracy: 0.5544 - loss:
1.3155
Epoch 55: val_accuracy did not improve from 0.48276
30/30 _____ 1s 28ms/step - accuracy: 0.5533 - loss:
1.3173 - val_accuracy: 0.4655 - val_loss: 1.4970
Epoch 56/100
30/30 _____ 0s 25ms/step - accuracy: 0.5232 - loss:
1.3861
Epoch 56: val_accuracy did not improve from 0.48276
30/30 _____ 1s 28ms/step - accuracy: 0.5229 - loss:
1.3856 - val_accuracy: 0.4655 - val_loss: 1.4840
Epoch 57/100
```

```
30/30 _____ 0s 25ms/step - accuracy: 0.5530 - loss: 1.3690
Epoch 57: val_accuracy did not improve from 0.48276
30/30 _____ 1s 27ms/step - accuracy: 0.5529 - loss: 1.3683 - val_accuracy: 0.4655 - val_loss: 1.4759
Epoch 57: early stopping
```

##Print the metrics

loss and accuracy graphs are plotted for all models in one graph later

```
from tensorflow.keras.models import load_model
model = load_model('best_model.keras')
print_scores(model, val_ds, train_ds)
```

```
Displaying accuracy and loss for Default
Test loss:1.4841
Test accuracy:0.4828
Train loss:1.3976
train accuracy:0.5247
```

##Modified CNN Model

```
model = Sequential([
    # Rescaling images
    layers.Rescaling(1./255, input_shape=(height, width, 3)),

    # First convolution layer.convolution with 32 filters with a 3x3
    kernel matrix
    layers.Conv2D(32, (3,3), padding='same', activation='relu'),
    layers.MaxPooling2D(), # Use max pooling
    layers.Dropout(0.1),

    # Second Convolution layer
    layers.Conv2D(32, (3,3), padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.1),

    # 3rd Convolution layer
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.1),

    # 4th Convolution layer
    layers.Conv2D(128, 3, padding='same', activation='relu'),
    layers.GlobalAveragePooling2D(),
    layers.Dropout(0.1),
```

```

#FCN
layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(128, activation='relu'),
layers.Dropout(0.2),
layers.Dense(len(class_names), activation='softmax')

])

```

##Compile the model

```

model.compile(optimizer='adam',

#https://www.tensorflow.org/api_docs/python/tf/keras/losses/CategoricalCrossentropy
#It states:
# "We expect labels to be provided in a one_hot
representation.
# If you want to provide labels as integers, please use
SparseCategoricalCrossentropy loss."
loss = 'sparse_categorical_crossentropy',
# keras has removed support for batch level precision
and recall metrics.Problem statement is obsolete
metrics=['accuracy']
)
model.summary()

```

Model: "sequential_1"

Layer (type) Param #	Output Shape
rescaling_1 (Rescaling) 0	(None, 224, 224, 3)
conv2d_2 (Conv2D) 896	(None, 224, 224, 32)
max_pooling2d_2 (MaxPooling2D) 0	(None, 112, 112, 32)
dropout (Dropout) 0	(None, 112, 112, 32)

conv2d_3 (Conv2D)	(None, 112, 112, 32)	
9,248		
max_pooling2d_3 (MaxPooling2D)	(None, 56, 56, 32)	
0		
dropout_1 (Dropout)	(None, 56, 56, 32)	
0		
conv2d_4 (Conv2D)	(None, 56, 56, 64)	
18,496		
max_pooling2d_4 (MaxPooling2D)	(None, 28, 28, 64)	
0		
dropout_2 (Dropout)	(None, 28, 28, 64)	
0		
conv2d_5 (Conv2D)	(None, 28, 28, 128)	
73,856		
global_average_pooling2d_1	(None, 128)	
0		
(GlobalAveragePooling2D)		
dropout_3 (Dropout)	(None, 128)	
0		
flatten_1 (Flatten)	(None, 128)	
0		
dense_3 (Dense)	(None, 128)	
16,512		
dense_4 (Dense)	(None, 128)	

16,512			
dropout_4 (Dropout)		(None, 128)	
0			
dense_5 (Dense)		(None, 9)	
1,161			

Total params: 136,681 (533.91 KB)

Trainable params: 136,681 (533.91 KB)

Non-trainable params: 0 (0.00 B)

##Train the model

```
histories = []

from tensorflow.keras.callbacks import EarlyStopping,ModelCheckpoint

early_stop=EarlyStopping(monitor='val_accuracy',mode='max',verbose=1,p
atience=10,min_delta=0.01)
checkpoint=ModelCheckpoint('best_model.keras',monitor='val_accuracy',m
ode='max',verbose=1,save_best_only=True)

epochs=100
history_c = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    batch_size = batch_size,
    callbacks = [early_stop,checkpoint]
)

histories.append(history_c)

Epoch 1/100
30/30 _____ 0s 178ms/step - accuracy: 0.3032 - loss:
2.0114
Epoch 1: val_accuracy improved from -inf to 0.30172, saving model to
best_model.keras
30/30 _____ 16s 242ms/step - accuracy: 0.3037 - loss:
2.0085 - val_accuracy: 0.3017 - val_loss: 1.8247
Epoch 2/100
30/30 _____ 0s 33ms/step - accuracy: 0.3535 - loss:
1.8137
```



```
Epoch 2: val_accuracy did not improve from 0.30172
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.3530 - loss:
1.8145 - val_accuracy: 0.3017 - val_loss: 1.8503
Epoch 3/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.3515 - loss:
1.8248
Epoch 3: val_accuracy did not improve from 0.30172
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.3511 - loss:
1.8244 - val_accuracy: 0.3017 - val_loss: 1.7972
Epoch 4/100
30/30 ━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.3525 - loss:
1.7807
Epoch 4: val_accuracy did not improve from 0.30172
30/30 ━━━━━━━━━━━ 1s 36ms/step - accuracy: 0.3522 - loss:
1.7817 - val_accuracy: 0.3017 - val_loss: 1.8086
Epoch 5/100
30/30 ━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.3359 - loss:
1.8012
Epoch 5: val_accuracy did not improve from 0.30172
30/30 ━━━━━━━━━━━ 1s 39ms/step - accuracy: 0.3361 - loss:
1.8012 - val_accuracy: 0.3017 - val_loss: 1.8066
Epoch 6/100
30/30 ━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.3538 - loss:
1.7908
Epoch 6: val_accuracy did not improve from 0.30172
30/30 ━━━━━━━━━━━ 1s 39ms/step - accuracy: 0.3534 - loss:
1.7909 - val_accuracy: 0.3017 - val_loss: 1.7699
Epoch 7/100
29/30 ━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.3090 - loss:
1.8624
Epoch 7: val_accuracy did not improve from 0.30172
30/30 ━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.3101 - loss:
1.8594 - val_accuracy: 0.3017 - val_loss: 1.7858
Epoch 8/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.3516 - loss:
1.7776
Epoch 8: val_accuracy improved from 0.30172 to 0.31897, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.3514 - loss:
1.7770 - val_accuracy: 0.3190 - val_loss: 1.7395
Epoch 9/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.3797 - loss:
1.7487
Epoch 9: val_accuracy improved from 0.31897 to 0.32328, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.3798 - loss:
1.7481 - val_accuracy: 0.3233 - val_loss: 1.7551
Epoch 10/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4095 - loss:
```

```
1.7087
Epoch 10: val_accuracy improved from 0.32328 to 0.33621, saving model
to best_model.keras
30/30 ━━━━━━━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.4089 - loss:
1.7095 - val_accuracy: 0.3362 - val_loss: 1.7165
Epoch 11/100
30/30 ━━━━━━━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4004 - loss:
1.6759
Epoch 11: val_accuracy did not improve from 0.33621
30/30 ━━━━━━━━━━━━━━━━━ 1s 36ms/step - accuracy: 0.3998 - loss:
1.6761 - val_accuracy: 0.3190 - val_loss: 1.6731
Epoch 12/100
29/30 ━━━━━━━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.3972 - loss:
1.6420
Epoch 12: val_accuracy did not improve from 0.33621
30/30 ━━━━━━━━━━━━━━━━━ 1s 36ms/step - accuracy: 0.3971 - loss:
1.6446 - val_accuracy: 0.3276 - val_loss: 1.7359
Epoch 13/100
30/30 ━━━━━━━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4046 - loss:
1.6948
Epoch 13: val_accuracy improved from 0.33621 to 0.35345, saving model
to best_model.keras
30/30 ━━━━━━━━━━━━━━━━━ 1s 40ms/step - accuracy: 0.4049 - loss:
1.6936 - val_accuracy: 0.3534 - val_loss: 1.6605
Epoch 14/100
30/30 ━━━━━━━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4402 - loss:
1.6305
Epoch 14: val_accuracy did not improve from 0.35345
30/30 ━━━━━━━━━━━━━━━━━ 1s 36ms/step - accuracy: 0.4405 - loss:
1.6287 - val_accuracy: 0.3534 - val_loss: 1.6316
Epoch 15/100
29/30 ━━━━━━━━━━━━━━━━━ 0s 35ms/step - accuracy: 0.4422 - loss:
1.6274
Epoch 15: val_accuracy improved from 0.35345 to 0.40948, saving model
to best_model.keras
30/30 ━━━━━━━━━━━━━━━━━ 1s 41ms/step - accuracy: 0.4418 - loss:
1.6249 - val_accuracy: 0.4095 - val_loss: 1.5927
Epoch 16/100
30/30 ━━━━━━━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.4701 - loss:
1.5659
Epoch 16: val_accuracy did not improve from 0.40948
30/30 ━━━━━━━━━━━━━━━━━ 1s 39ms/step - accuracy: 0.4700 - loss:
1.5657 - val_accuracy: 0.3922 - val_loss: 1.5756
Epoch 17/100
30/30 ━━━━━━━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.4523 - loss:
1.5171
Epoch 17: val_accuracy improved from 0.40948 to 0.43103, saving model
to best_model.keras
30/30 ━━━━━━━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.4523 - loss:
```

```
1.5182 - val_accuracy: 0.4310 - val_loss: 1.5797
Epoch 18/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4745 - loss:
1.4854
Epoch 18: val_accuracy did not improve from 0.43103
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.4746 - loss:
1.4863 - val_accuracy: 0.4267 - val_loss: 1.5424
Epoch 19/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.5131 - loss:
1.4194
Epoch 19: val_accuracy improved from 0.43103 to 0.45259, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.5125 - loss:
1.4218 - val_accuracy: 0.4526 - val_loss: 1.5570
Epoch 20/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.5114 - loss:
1.4332
Epoch 20: val_accuracy did not improve from 0.45259
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.5108 - loss:
1.4347 - val_accuracy: 0.4224 - val_loss: 1.5428
Epoch 21/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4983 - loss:
1.4524
Epoch 21: val_accuracy did not improve from 0.45259
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.4983 - loss:
1.4529 - val_accuracy: 0.3922 - val_loss: 1.5954
Epoch 22/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4794 - loss:
1.4627
Epoch 22: val_accuracy did not improve from 0.45259
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.4798 - loss:
1.4627 - val_accuracy: 0.4267 - val_loss: 1.5085
Epoch 23/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.5051 - loss:
1.4491
Epoch 23: val_accuracy improved from 0.45259 to 0.47845, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 40ms/step - accuracy: 0.5049 - loss:
1.4500 - val_accuracy: 0.4784 - val_loss: 1.5350
Epoch 24/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.4906 - loss:
1.4845
Epoch 24: val_accuracy did not improve from 0.47845
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.4908 - loss:
1.4839 - val_accuracy: 0.4353 - val_loss: 1.4958
Epoch 25/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.5208 - loss:
1.4296
Epoch 25: val_accuracy did not improve from 0.47845
```

```
30/30 _____ 1s 35ms/step - accuracy: 0.5212 - loss:
1.4287 - val_accuracy: 0.4698 - val_loss: 1.5100
Epoch 26/100
30/30 _____ 0s 34ms/step - accuracy: 0.5185 - loss:
1.4756
Epoch 26: val_accuracy improved from 0.47845 to 0.50431, saving model
to best_model.keras
30/30 _____ 1s 41ms/step - accuracy: 0.5187 - loss:
1.4736 - val_accuracy: 0.5043 - val_loss: 1.4486
Epoch 27/100
30/30 _____ 0s 34ms/step - accuracy: 0.5326 - loss:
1.3288
Epoch 27: val_accuracy did not improve from 0.50431
30/30 _____ 1s 39ms/step - accuracy: 0.5326 - loss:
1.3305 - val_accuracy: 0.4741 - val_loss: 1.4511
Epoch 28/100
30/30 _____ 0s 33ms/step - accuracy: 0.5674 - loss:
1.3374
Epoch 28: val_accuracy did not improve from 0.50431
30/30 _____ 1s 36ms/step - accuracy: 0.5662 - loss:
1.3401 - val_accuracy: 0.4612 - val_loss: 1.4757
Epoch 29/100
30/30 _____ 0s 33ms/step - accuracy: 0.5563 - loss:
1.3446
Epoch 29: val_accuracy improved from 0.50431 to 0.55603, saving model
to best_model.keras
30/30 _____ 1s 37ms/step - accuracy: 0.5562 - loss:
1.3455 - val_accuracy: 0.5560 - val_loss: 1.4172
Epoch 30/100
30/30 _____ 0s 33ms/step - accuracy: 0.5409 - loss:
1.3657
Epoch 30: val_accuracy improved from 0.55603 to 0.58190, saving model
to best_model.keras
30/30 _____ 1s 37ms/step - accuracy: 0.5409 - loss:
1.3661 - val_accuracy: 0.5819 - val_loss: 1.4059
Epoch 31/100
30/30 _____ 0s 32ms/step - accuracy: 0.5608 - loss:
1.3830
Epoch 31: val_accuracy did not improve from 0.58190
30/30 _____ 1s 35ms/step - accuracy: 0.5608 - loss:
1.3820 - val_accuracy: 0.5172 - val_loss: 1.4210
Epoch 32/100
30/30 _____ 0s 33ms/step - accuracy: 0.5864 - loss:
1.2669
Epoch 32: val_accuracy did not improve from 0.58190
30/30 _____ 1s 35ms/step - accuracy: 0.5866 - loss:
1.2666 - val_accuracy: 0.5517 - val_loss: 1.4315
Epoch 33/100
29/30 _____ 0s 33ms/step - accuracy: 0.5997 - loss:
```

```
1.2963
Epoch 33: val_accuracy did not improve from 0.58190
30/30 ━━━━━━━━━━━ 1s 39ms/step - accuracy: 0.5997 - loss:
1.2944 - val_accuracy: 0.5603 - val_loss: 1.3752
Epoch 34/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6008 - loss:
1.2259
Epoch 34: val_accuracy did not improve from 0.58190
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.6007 - loss:
1.2260 - val_accuracy: 0.5086 - val_loss: 1.4549
Epoch 35/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.5908 - loss:
1.2610
Epoch 35: val_accuracy did not improve from 0.58190
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.5900 - loss:
1.2624 - val_accuracy: 0.5776 - val_loss: 1.3698
Epoch 36/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6070 - loss:
1.2713
Epoch 36: val_accuracy improved from 0.58190 to 0.61207, saving model
to best_model.keras
30/30 ━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.6068 - loss:
1.2711 - val_accuracy: 0.6121 - val_loss: 1.2901
Epoch 37/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.5647 - loss:
1.2750
Epoch 37: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 39ms/step - accuracy: 0.5655 - loss:
1.2737 - val_accuracy: 0.5819 - val_loss: 1.3531
Epoch 38/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6387 - loss:
1.2129
Epoch 38: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 39ms/step - accuracy: 0.6383 - loss:
1.2117 - val_accuracy: 0.5991 - val_loss: 1.3431
Epoch 39/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6625 - loss:
1.1206
Epoch 39: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 36ms/step - accuracy: 0.6619 - loss:
1.1210 - val_accuracy: 0.5560 - val_loss: 1.4739
Epoch 40/100
29/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6515 - loss:
1.1457
Epoch 40: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.6503 - loss:
1.1476 - val_accuracy: 0.5819 - val_loss: 1.2794
Epoch 41/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6806 - loss:
```

```

1.0863
Epoch 41: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.6798 - loss:
1.0882 - val_accuracy: 0.5991 - val_loss: 1.3043
Epoch 42/100
30/30 ━━━━━━━━━━━ 0s 32ms/step - accuracy: 0.6728 - loss:
1.0701
Epoch 42: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.6722 - loss:
1.0718 - val_accuracy: 0.6034 - val_loss: 1.3370
Epoch 43/100
30/30 ━━━━━━━━━━━ 0s 32ms/step - accuracy: 0.6360 - loss:
1.1083
Epoch 43: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 37ms/step - accuracy: 0.6363 - loss:
1.1085 - val_accuracy: 0.6034 - val_loss: 1.2843
Epoch 44/100
30/30 ━━━━━━━━━━━ 0s 33ms/step - accuracy: 0.6622 - loss:
1.0919
Epoch 44: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 38ms/step - accuracy: 0.6617 - loss:
1.0917 - val_accuracy: 0.5690 - val_loss: 1.3548
Epoch 45/100
29/30 ━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.6657 - loss:
1.0585
Epoch 45: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 35ms/step - accuracy: 0.6655 - loss:
1.0596 - val_accuracy: 0.5905 - val_loss: 1.2406
Epoch 46/100
29/30 ━━━━━━━━━━━ 0s 34ms/step - accuracy: 0.6441 - loss:
1.1389
Epoch 46: val_accuracy did not improve from 0.61207
30/30 ━━━━━━━━━━━ 1s 36ms/step - accuracy: 0.6436 - loss:
1.1382 - val_accuracy: 0.6078 - val_loss: 1.2634
Epoch 46: early stopping

```

##Print Metrics

```

from tensorflow.keras.models import load_model
model = load_model('best_model.keras')
print_scores(model, val_ds, train_ds, 'CNN Model')

```

```


Displaying accuracy and loss for CNN Model
Test loss:1.2901
Test accuracy:0.6121
Train loss:1.1262
train accuracy:0.6409

```

##Display the predictions

```
image_map = get_predicted_labels(model,class_names)
display_test_images(image_map,'CNN Model. Title is the predicted
label')
```

Predictions for CNN Model. Title is the predicted label



Notice how classes with very few images in training set tend to get misclassified. sailboat , which has the largest number of images is the best predicted class

##Confusion Matrix

##Get the predicted and True labels


```
true_labels, predicted_labels = get_metrics(model, val_ds)
```

```
##Display Confusion Matrix
```


```
display_cm(true_labels, predicted_labels)
```

Confusion Matrix

```
[[ 0  2  0  0  0  0  8  0  3]
 [ 0 33  0  0  3  0  5  0  4]
 [ 0  9  0  0  1  0  2  0  3]
 [ 0  0  0  0  0  0  1  0  2]
 [ 0  1  0  0 25  0  3  0  4]
 [ 0  0  0  0  0  0  2  0  0]
 [ 0  3  0  0  3  0 30  0  7]
 [ 0  0  0  0  0  0  2  0  1]
 [ 0  8  0  0  4  0  8  0 47]]
```

Confusion Matrix Displayed

<Figure size 1000x1000 with 0 Axes>



As expected Sailboat,Kayak,Gondola and Cruise ship has high true positive values because of large number of images in the training data. Most mis-classifications are seen in the other image classes

##Classification Report

```
cl_report = classification_report(true_labels,predicted_labels)
print(cl_report)

print('Index to name mapping:\n')

for k,v in class_dict.items():
    print(f'{k} - > {v}')
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	13
1	0.59	0.73	0.65	45
2	0.00	0.00	0.00	15

3	0.00	0.00	0.00	3
4	0.69	0.76	0.72	33
5	0.00	0.00	0.00	2
6	0.49	0.70	0.58	43
7	0.00	0.00	0.00	3
8	0.66	0.70	0.68	67
accuracy			0.60	224
macro avg	0.27	0.32	0.29	224
weighted avg	0.51	0.60	0.55	224

Index to name mapping:

```
0 - > inflatable_boat
1 - > freight_boat
2 - > paper_boat
3 - > buoy
4 - > ferry_boat
5 - > cruise_ship
6 - > gondola
7 - > kayak
8 - > sailboat
```

As expected the classes with higher support have the better overall scores

#Transfer Learning with MobileNet application

##Import mobilenet_v2 and freeze the weights

```
from tensorflow.keras.applications.mobilenet_v2 import MobileNetV2
mob_model = tf.keras.applications.MobileNet(
    input_shape=(height,width,3),
    include_top=False,
    weights="imagenet",
    input_tensor=None,
    pooling=None,
    classes=1000,
    classifier_activation="softmax",
    name='MobileNet',
)
mob_model.trainable=False
mob_model.summary()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/mobilenet/mobilenet_1_0_224_tf_no_top.h5
17225924/17225924 ————— 0s 0us/step
```

```
Model: "MobileNet"
```

Layer (type) Param #	Output Shape
input_layer_2 (InputLayer) 0	(None, 224, 224, 3)
conv1 (Conv2D) 864	(None, 112, 112, 32)
conv1_bn (BatchNormalization) 128	(None, 112, 112, 32)
conv1_relu (ReLU) 0	(None, 112, 112, 32)
conv_dw_1 (DepthwiseConv2D) 288	(None, 112, 112, 32)
conv_dw_1_bn (BatchNormalization) 128	(None, 112, 112, 32)
conv_dw_1_relu (ReLU) 0	(None, 112, 112, 32)
conv_pw_1 (Conv2D) 2,048	(None, 112, 112, 64)
conv_pw_1_bn (BatchNormalization) 256	(None, 112, 112, 64)
conv_pw_1_relu (ReLU) 0	(None, 112, 112, 64)
conv_pad_2 (ZeroPadding2D) 0	(None, 113, 113, 64)

576	conv_dw_2 (DepthwiseConv2D)	(None, 56, 56, 64)	
256	conv_dw_2_bn (BatchNormalization)	(None, 56, 56, 64)	
0	conv_dw_2_relu (ReLU)	(None, 56, 56, 64)	
8,192	conv_pw_2 (Conv2D)	(None, 56, 56, 128)	
512	conv_pw_2_bn (BatchNormalization)	(None, 56, 56, 128)	
0	conv_pw_2_relu (ReLU)	(None, 56, 56, 128)	
1,152	conv_dw_3 (DepthwiseConv2D)	(None, 56, 56, 128)	
512	conv_dw_3_bn (BatchNormalization)	(None, 56, 56, 128)	
0	conv_dw_3_relu (ReLU)	(None, 56, 56, 128)	
16,384	conv_pw_3 (Conv2D)	(None, 56, 56, 128)	
512	conv_pw_3_bn (BatchNormalization)	(None, 56, 56, 128)	
0	conv_pw_3_relu (ReLU)	(None, 56, 56, 128)	
0	conv_pad_4 (ZeroPadding2D)	(None, 57, 57, 128)	

conv_dw_4 (DepthwiseConv2D)	(None, 28, 28, 128)	
1,152		
conv_dw_4_bn (BatchNormalization)	(None, 28, 28, 128)	
512		
conv_dw_4_relu (ReLU)	(None, 28, 28, 128)	
0		
conv_pw_4 (Conv2D)	(None, 28, 28, 256)	
32,768		
conv_pw_4_bn (BatchNormalization)	(None, 28, 28, 256)	
1,024		
conv_pw_4_relu (ReLU)	(None, 28, 28, 256)	
0		
conv_dw_5 (DepthwiseConv2D)	(None, 28, 28, 256)	
2,304		
conv_dw_5_bn (BatchNormalization)	(None, 28, 28, 256)	
1,024		
conv_dw_5_relu (ReLU)	(None, 28, 28, 256)	
0		
conv_pw_5 (Conv2D)	(None, 28, 28, 256)	
65,536		
conv_pw_5_bn (BatchNormalization)	(None, 28, 28, 256)	
1,024		
conv_pw_5_relu (ReLU)	(None, 28, 28, 256)	
0		

0	conv_pad_6 (ZeroPadding2D)	(None, 29, 29, 256)	
2,304	conv_dw_6 (DepthwiseConv2D)	(None, 14, 14, 256)	
1,024	conv_dw_6_bn (BatchNormalization)	(None, 14, 14, 256)	
0	conv_dw_6_relu (ReLU)	(None, 14, 14, 256)	
131,072	conv_pw_6 (Conv2D)	(None, 14, 14, 512)	
2,048	conv_pw_6_bn (BatchNormalization)	(None, 14, 14, 512)	
0	conv_pw_6_relu (ReLU)	(None, 14, 14, 512)	
4,608	conv_dw_7 (DepthwiseConv2D)	(None, 14, 14, 512)	
2,048	conv_dw_7_bn (BatchNormalization)	(None, 14, 14, 512)	
0	conv_dw_7_relu (ReLU)	(None, 14, 14, 512)	
262,144	conv_pw_7 (Conv2D)	(None, 14, 14, 512)	
2,048	conv_pw_7_bn (BatchNormalization)	(None, 14, 14, 512)	
0	conv_pw_7_relu (ReLU)	(None, 14, 14, 512)	

4,608	conv_dw_8 (DepthwiseConv2D)	(None, 14, 14, 512)
2,048	conv_dw_8_bn (BatchNormalization)	(None, 14, 14, 512)
0	conv_dw_8_relu (ReLU)	(None, 14, 14, 512)
262,144	conv_pw_8 (Conv2D)	(None, 14, 14, 512)
2,048	conv_pw_8_bn (BatchNormalization)	(None, 14, 14, 512)
0	conv_pw_8_relu (ReLU)	(None, 14, 14, 512)
4,608	conv_dw_9 (DepthwiseConv2D)	(None, 14, 14, 512)
2,048	conv_dw_9_bn (BatchNormalization)	(None, 14, 14, 512)
0	conv_dw_9_relu (ReLU)	(None, 14, 14, 512)
262,144	conv_pw_9 (Conv2D)	(None, 14, 14, 512)
2,048	conv_pw_9_bn (BatchNormalization)	(None, 14, 14, 512)
0	conv_pw_9_relu (ReLU)	(None, 14, 14, 512)

conv_dw_10 (DepthwiseConv2D)	(None, 14, 14, 512)
4,608	
conv_dw_10_bn (BatchNormalization)	(None, 14, 14, 512)
2,048	
conv_dw_10_relu (ReLU)	(None, 14, 14, 512)
0	
conv_pw_10 (Conv2D)	(None, 14, 14, 512)
262,144	
conv_pw_10_bn (BatchNormalization)	(None, 14, 14, 512)
2,048	
conv_pw_10_relu (ReLU)	(None, 14, 14, 512)
0	
conv_dw_11 (DepthwiseConv2D)	(None, 14, 14, 512)
4,608	
conv_dw_11_bn (BatchNormalization)	(None, 14, 14, 512)
2,048	
conv_dw_11_relu (ReLU)	(None, 14, 14, 512)
0	
conv_pw_11 (Conv2D)	(None, 14, 14, 512)
262,144	
conv_pw_11_bn (BatchNormalization)	(None, 14, 14, 512)
2,048	
conv_pw_11_relu (ReLU)	(None, 14, 14, 512)
0	
conv_pad_12 (ZeroPadding2D)	(None, 15, 15, 512)
0	

conv_dw_12 (DepthwiseConv2D)	(None, 7, 7, 512)	
4,608		
conv_dw_12_bn (BatchNormalization)	(None, 7, 7, 512)	
2,048		
conv_dw_12_relu (ReLU)	(None, 7, 7, 512)	
0		
conv_pw_12 (Conv2D)	(None, 7, 7, 1024)	
524,288		
conv_pw_12_bn (BatchNormalization)	(None, 7, 7, 1024)	
4,096		
conv_pw_12_relu (ReLU)	(None, 7, 7, 1024)	
0		
conv_dw_13 (DepthwiseConv2D)	(None, 7, 7, 1024)	
9,216		
conv_dw_13_bn (BatchNormalization)	(None, 7, 7, 1024)	
4,096		
conv_dw_13_relu (ReLU)	(None, 7, 7, 1024)	
0		
conv_pw_13 (Conv2D)	(None, 7, 7, 1024)	
1,048,576		
conv_pw_13_bn (BatchNormalization)	(None, 7, 7, 1024)	
4,096		
conv_pw_13_relu (ReLU)	(None, 7, 7, 1024)	
0		

```
Total params: 3,228,864 (12.32 MB)
Trainable params: 0 (0.00 B)
Non-trainable params: 3,228,864 (12.32 MB)
```

```
##Build an FCN
```

```
##Get training and validation set at 70:30 ratio and seed 1
```

```
train_ds = get_ds(data_dir,validation_split=0.3,subset='training',seed
= 1,image_size=(height,width),batch_size=batch_size)
val_ds = get_ds(data_dir,validation_split=0.3,subset='validation',seed
= 1,image_size=(height,width),batch_size=batch_size)
```

```
Found 1162 files belonging to 9 classes.
Using 814 files for training.
Found 1162 files belonging to 9 classes.
Using 348 files for validation.
```

```
##Set up dataset with caching and prefetch
```

```
AUTOTUNE = tf.data.AUTOTUNE

train_ds =
train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```
##Add a FCN to the mobile net model(no data augmentation)
```

```
model_t = create_transfer_model(mob_model)
model_t.name = 'Transfer_Model'
```

```
Model: "sequential_2"
```

Layer (type) Param #	Output Shape
rescaling_2 (Rescaling) 0	(None, 224, 224, 3)
MobileNet (Functional)	(None, 7, 7, 1024)

3,228,864		
0	global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 1024)
0	dropout_5 (Dropout)	(None, 1024)
0	flatten_2 (Flatten)	(None, 1024)
262,400	dense_6 (Dense)	(None, 256)
1,024	batch_normalization (BatchNormalization)	(None, 256)
0	dropout_6 (Dropout)	(None, 256)
32,896	dense_7 (Dense)	(None, 128)
512	batch_normalization_1 (BatchNormalization)	(None, 128)
0	dropout_7 (Dropout)	(None, 128)
1,161	dense_8 (Dense)	(None, 9)

```
Total params: 3,526,857 (13.45 MB)
Trainable params: 297,225 (1.13 MB)
Non-trainable params: 3,229,632 (12.32 MB)
```

```
##Compile the model
```

```
model_t.compile(optimizer='adam',
loss='sparse_categorical_crossentropy', metrics=['accuracy'])
print('model_t compiled')

model_t compiled
```

```
##Configure for early stopping and Checkpoint and train the model
```

```
early_stop=EarlyStopping(monitor='val_accuracy',mode='max',verbose=1,p
atience=10,min_delta=0.01)
checkpoint=ModelCheckpoint('best_model_t.keras',monitor='val_accuracy'
,mode='max',verbose=1,save_best_only=True)

epochs=100
history_t = model_t.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    batch_size = batch_size,
    callbacks = [early_stop,checkpoint]
)

histories.append(history_t)

Epoch 1/100
26/26 _____ 0s 220ms/step - accuracy: 0.4264 - loss:
1.8711
Epoch 1: val_accuracy improved from -inf to 0.67529, saving model to
best_model_t.keras
26/26 _____ 28s 513ms/step - accuracy: 0.4328 - loss:
1.8516 - val_accuracy: 0.6753 - val_loss: 1.0619
Epoch 2/100
25/26 _____ 0s 22ms/step - accuracy: 0.8425 - loss:
0.5144
Epoch 2: val_accuracy improved from 0.67529 to 0.77874, saving model
to best_model_t.keras
26/26 _____ 1s 43ms/step - accuracy: 0.8447 - loss:
0.5084 - val_accuracy: 0.7787 - val_loss: 0.7892
Epoch 3/100
24/26 _____ 0s 22ms/step - accuracy: 0.9499 - loss:
0.2008
```

Epoch 3: val_accuracy improved from 0.77874 to 0.80172, saving model to best_model_t.keras
26/26 _____ 1s 43ms/step - accuracy: 0.9506 - loss: 0.2012 - val_accuracy: 0.8017 - val_loss: 0.6515
Epoch 4/100
25/26 _____ 0s 22ms/step - accuracy: 0.9720 - loss: 0.1307
Epoch 4: val_accuracy improved from 0.80172 to 0.81897, saving model to best_model_t.keras
26/26 _____ 1s 52ms/step - accuracy: 0.9725 - loss: 0.1305 - val_accuracy: 0.8190 - val_loss: 0.5712
Epoch 5/100
24/26 _____ 0s 23ms/step - accuracy: 0.9926 - loss: 0.0737
Epoch 5: val_accuracy improved from 0.81897 to 0.83046, saving model to best_model_t.keras
26/26 _____ 2s 44ms/step - accuracy: 0.9921 - loss: 0.0745 - val_accuracy: 0.8305 - val_loss: 0.5724
Epoch 6/100
25/26 _____ 0s 22ms/step - accuracy: 0.9859 - loss: 0.0759
Epoch 6: val_accuracy improved from 0.83046 to 0.85057, saving model to best_model_t.keras
26/26 _____ 1s 44ms/step - accuracy: 0.9863 - loss: 0.0754 - val_accuracy: 0.8506 - val_loss: 0.5604
Epoch 7/100
25/26 _____ 0s 22ms/step - accuracy: 0.9990 - loss: 0.0547
Epoch 7: val_accuracy did not improve from 0.85057
26/26 _____ 1s 32ms/step - accuracy: 0.9988 - loss: 0.0545 - val_accuracy: 0.8391 - val_loss: 0.5507
Epoch 8/100
25/26 _____ 0s 22ms/step - accuracy: 1.0000 - loss: 0.0334
Epoch 8: val_accuracy did not improve from 0.85057
26/26 _____ 1s 34ms/step - accuracy: 1.0000 - loss: 0.0336 - val_accuracy: 0.8506 - val_loss: 0.5217
Epoch 9/100
25/26 _____ 0s 22ms/step - accuracy: 0.9998 - loss: 0.0245
Epoch 9: val_accuracy did not improve from 0.85057
26/26 _____ 1s 32ms/step - accuracy: 0.9997 - loss: 0.0247 - val_accuracy: 0.8506 - val_loss: 0.5314
Epoch 10/100
25/26 _____ 0s 22ms/step - accuracy: 0.9995 - loss: 0.0188
Epoch 10: val_accuracy did not improve from 0.85057
26/26 _____ 1s 32ms/step - accuracy: 0.9993 - loss: 0.0192 - val_accuracy: 0.8477 - val_loss: 0.5483

```
Epoch 11/100
25/26 _____ 0s 22ms/step - accuracy: 0.9997 - loss:
0.0218
Epoch 11: val_accuracy did not improve from 0.85057
26/26 _____ 1s 32ms/step - accuracy: 0.9994 - loss:
0.0222 - val_accuracy: 0.8420 - val_loss: 0.5499
Epoch 12/100
25/26 _____ 0s 22ms/step - accuracy: 0.9951 - loss:
0.0248
Epoch 12: val_accuracy improved from 0.85057 to 0.86207, saving model
to best_model_t.keras
26/26 _____ 1s 43ms/step - accuracy: 0.9952 - loss:
0.0249 - val_accuracy: 0.8621 - val_loss: 0.5267
Epoch 13/100
25/26 _____ 0s 22ms/step - accuracy: 1.0000 - loss:
0.0149
Epoch 13: val_accuracy did not improve from 0.86207
26/26 _____ 1s 32ms/step - accuracy: 1.0000 - loss:
0.0150 - val_accuracy: 0.8506 - val_loss: 0.5698
Epoch 14/100
25/26 _____ 0s 23ms/step - accuracy: 1.0000 - loss:
0.0154
Epoch 14: val_accuracy did not improve from 0.86207
26/26 _____ 1s 32ms/step - accuracy: 1.0000 - loss:
0.0154 - val_accuracy: 0.8621 - val_loss: 0.5414
Epoch 15/100
25/26 _____ 0s 23ms/step - accuracy: 1.0000 - loss:
0.0099
Epoch 15: val_accuracy did not improve from 0.86207
26/26 _____ 1s 36ms/step - accuracy: 1.0000 - loss:
0.0100 - val_accuracy: 0.8534 - val_loss: 0.5469
Epoch 16/100
25/26 _____ 0s 23ms/step - accuracy: 0.9983 - loss:
0.0114
Epoch 16: val_accuracy did not improve from 0.86207
26/26 _____ 1s 33ms/step - accuracy: 0.9983 - loss:
0.0114 - val_accuracy: 0.8506 - val_loss: 0.5624
Epoch 17/100
25/26 _____ 0s 23ms/step - accuracy: 1.0000 - loss:
0.0088
Epoch 17: val_accuracy did not improve from 0.86207
26/26 _____ 1s 36ms/step - accuracy: 1.0000 - loss:
0.0089 - val_accuracy: 0.8563 - val_loss: 0.5694
Epoch 18/100
24/26 _____ 0s 22ms/step - accuracy: 1.0000 - loss:
0.0099
Epoch 18: val_accuracy did not improve from 0.86207
26/26 _____ 1s 32ms/step - accuracy: 1.0000 - loss:
0.0097 - val_accuracy: 0.8362 - val_loss: 0.6015
```

```

Epoch 19/100
25/26 _____ 0s 23ms/step - accuracy: 0.9988 - loss:
0.0097
Epoch 19: val_accuracy did not improve from 0.86207
26/26 _____ 1s 33ms/step - accuracy: 0.9988 - loss:
0.0097 - val_accuracy: 0.8362 - val_loss: 0.5923
Epoch 20/100
25/26 _____ 0s 22ms/step - accuracy: 0.9975 - loss:
0.0217
Epoch 20: val_accuracy did not improve from 0.86207
26/26 _____ 1s 32ms/step - accuracy: 0.9976 - loss:
0.0212 - val_accuracy: 0.8477 - val_loss: 0.5721
Epoch 21/100
25/26 _____ 0s 22ms/step - accuracy: 1.0000 - loss:
0.0094
Epoch 21: val_accuracy did not improve from 0.86207
26/26 _____ 1s 32ms/step - accuracy: 1.0000 - loss:
0.0093 - val_accuracy: 0.8534 - val_loss: 0.5480
Epoch 22/100
25/26 _____ 0s 22ms/step - accuracy: 0.9971 - loss:
0.0094
Epoch 22: val_accuracy did not improve from 0.86207
26/26 _____ 1s 34ms/step - accuracy: 0.9972 - loss:
0.0092 - val_accuracy: 0.8563 - val_loss: 0.5547
Epoch 22: early stopping

```

##Print metrics

```

from tensorflow.keras.models import load_model
model_t = load_model('best_model_t.keras')
print_scores(model_t, val_ds, train_ds, 'Transfer Model')

Displaying accuracy and loss for Transfer Model
Test loss:0.5267
Test accuracy:0.8621
Train loss:0.0102
train accuracy:0.9988

```

Severe overfitting can be observed. We will modify this model by increasing dropouts. The below model uses a dropout that i have arrived at after trial and error which produced lesser overfitting

##Dropout adjusted Transfer Model(no augmentation)

```

model_t1 = create_transfer_model(mob_model, dropout=0.4)
Model: "sequential_3"

```

Layer (type) Param #	Output Shape
rescaling_3 (Rescaling) 0	(None, 224, 224, 3)
MobileNet (Functional) 3,228,864	(None, 7, 7, 1024)
global_average_pooling2d_3 0 (GlobalAveragePooling2D)	(None, 1024)
dropout_8 (Dropout) 0	(None, 1024)
flatten_3 (Flatten) 0	(None, 1024)
dense_9 (Dense) 262,400	(None, 256)
batch_normalization_2 1,024 (BatchNormalization)	(None, 256)
dropout_9 (Dropout) 0	(None, 256)
dense_10 (Dense) 32,896	(None, 128)
batch_normalization_3 512 (BatchNormalization)	(None, 128)

0	dropout_10 (Dropout)	(None, 128)
1,161	dense_11 (Dense)	(None, 9)

Total params: 3,526,857 (13.45 MB)

Trainable params: 297,225 (1.13 MB)

Non-trainable params: 3,229,632 (12.32 MB)

##Compile the model

```
model_t1.compile(optimizer='adam',
loss='sparse_categorical_crossentropy', metrics=['accuracy'])
print('model_t1 compiled')
```

model_t1 compiled

##Train the model

```
early_stop=EarlyStopping(monitor='val_accuracy',mode='max',verbose=1,p
atience=10,min_delta=0.01)
checkpoint=ModelCheckpoint('best_model_t1.keras',monitor='val_accuracy
',mode='max',verbose=1,save_best_only=True)
```

```
epochs=100
history_t1 = model_t1.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    batch_size = batch_size,
    callbacks = [early_stop,checkpoint]
)
```

histories.append(history_t1)

Epoch 1/100

25/26 ————— 0s 149ms/step - accuracy: 0.2050 - loss: 2.7388

Epoch 1: val_accuracy improved from -inf to 0.52874, saving model to best_model_t1.keras

26/26 ————— 15s 313ms/step - accuracy: 0.2133 - loss: 2.7025 - val_accuracy: 0.5287 - val_loss: 1.4464

```
Epoch 2/100
25/26 _____ 0s 22ms/step - accuracy: 0.5586 - loss:
1.4237
Epoch 2: val_accuracy improved from 0.52874 to 0.70977, saving model
to best_model_t1.keras
26/26 _____ 1s 43ms/step - accuracy: 0.5644 - loss:
1.4091 - val_accuracy: 0.7098 - val_loss: 1.0327
Epoch 3/100
25/26 _____ 0s 22ms/step - accuracy: 0.7462 - loss:
0.8408
Epoch 3: val_accuracy improved from 0.70977 to 0.76724, saving model
to best_model_t1.keras
26/26 _____ 1s 43ms/step - accuracy: 0.7441 - loss:
0.8477 - val_accuracy: 0.7672 - val_loss: 0.8145
Epoch 4/100
25/26 _____ 0s 22ms/step - accuracy: 0.8037 - loss:
0.6875
Epoch 4: val_accuracy improved from 0.76724 to 0.80172, saving model
to best_model_t1.keras
26/26 _____ 1s 44ms/step - accuracy: 0.8027 - loss:
0.6895 - val_accuracy: 0.8017 - val_loss: 0.6673
Epoch 5/100
25/26 _____ 0s 22ms/step - accuracy: 0.8155 - loss:
0.6046
Epoch 5: val_accuracy did not improve from 0.80172
26/26 _____ 1s 32ms/step - accuracy: 0.8143 - loss:
0.6059 - val_accuracy: 0.7902 - val_loss: 0.6345
Epoch 6/100
25/26 _____ 0s 22ms/step - accuracy: 0.7889 - loss:
0.6165
Epoch 6: val_accuracy improved from 0.80172 to 0.82184, saving model
to best_model_t1.keras
26/26 _____ 1s 43ms/step - accuracy: 0.7899 - loss:
0.6152 - val_accuracy: 0.8218 - val_loss: 0.6081
Epoch 7/100
25/26 _____ 0s 22ms/step - accuracy: 0.8097 - loss:
0.6110
Epoch 7: val_accuracy improved from 0.82184 to 0.83046, saving model
to best_model_t1.keras
26/26 _____ 1s 46ms/step - accuracy: 0.8105 - loss:
0.6066 - val_accuracy: 0.8305 - val_loss: 0.5926
Epoch 8/100
25/26 _____ 0s 22ms/step - accuracy: 0.8504 - loss:
0.4538
Epoch 8: val_accuracy did not improve from 0.83046
26/26 _____ 1s 32ms/step - accuracy: 0.8499 - loss:
0.4551 - val_accuracy: 0.8276 - val_loss: 0.5940
Epoch 9/100
25/26 _____ 0s 22ms/step - accuracy: 0.8656 - loss:
```

```
0.4162
Epoch 9: val_accuracy did not improve from 0.83046
26/26 _____ 1s 33ms/step - accuracy: 0.8653 - loss:
0.4172 - val_accuracy: 0.8276 - val_loss: 0.5682
Epoch 10/100
24/26 _____ 0s 24ms/step - accuracy: 0.8694 - loss:
0.3948
Epoch 10: val_accuracy improved from 0.83046 to 0.83908, saving model
to best_model_t1.keras
26/26 _____ 1s 51ms/step - accuracy: 0.8682 - loss:
0.3973 - val_accuracy: 0.8391 - val_loss: 0.5561
Epoch 11/100
25/26 _____ 0s 22ms/step - accuracy: 0.8981 - loss:
0.3393
Epoch 11: val_accuracy did not improve from 0.83908
26/26 _____ 2s 32ms/step - accuracy: 0.8967 - loss:
0.3436 - val_accuracy: 0.8362 - val_loss: 0.5802
Epoch 12/100
25/26 _____ 0s 23ms/step - accuracy: 0.8543 - loss:
0.4477
Epoch 12: val_accuracy did not improve from 0.83908
26/26 _____ 1s 33ms/step - accuracy: 0.8558 - loss:
0.4441 - val_accuracy: 0.8391 - val_loss: 0.5603
Epoch 13/100
24/26 _____ 0s 23ms/step - accuracy: 0.8836 - loss:
0.3455
Epoch 13: val_accuracy improved from 0.83908 to 0.84483, saving model
to best_model_t1.keras
26/26 _____ 2s 44ms/step - accuracy: 0.8836 - loss:
0.3455 - val_accuracy: 0.8448 - val_loss: 0.5482
Epoch 14/100
25/26 _____ 0s 22ms/step - accuracy: 0.8809 - loss:
0.3696
Epoch 14: val_accuracy did not improve from 0.84483
26/26 _____ 1s 32ms/step - accuracy: 0.8827 - loss:
0.3661 - val_accuracy: 0.8362 - val_loss: 0.5369
Epoch 15/100
25/26 _____ 0s 22ms/step - accuracy: 0.8954 - loss:
0.3131
Epoch 15: val_accuracy did not improve from 0.84483
26/26 _____ 1s 32ms/step - accuracy: 0.8961 - loss:
0.3114 - val_accuracy: 0.8391 - val_loss: 0.5517
Epoch 16/100
25/26 _____ 0s 22ms/step - accuracy: 0.9285 - loss:
0.2531
Epoch 16: val_accuracy did not improve from 0.84483
26/26 _____ 1s 32ms/step - accuracy: 0.9278 - loss:
0.2533 - val_accuracy: 0.8391 - val_loss: 0.5259
Epoch 17/100
```

```
25/26 _____ 0s 22ms/step - accuracy: 0.9241 - loss:
0.2381
Epoch 17: val_accuracy did not improve from 0.84483
26/26 _____ 1s 32ms/step - accuracy: 0.9240 - loss:
0.2376 - val_accuracy: 0.8276 - val_loss: 0.5429
Epoch 18/100
25/26 _____ 0s 22ms/step - accuracy: 0.9418 - loss:
0.2094
Epoch 18: val_accuracy improved from 0.84483 to 0.84770, saving model
to best_model_t1.keras
26/26 _____ 2s 46ms/step - accuracy: 0.9413 - loss:
0.2100 - val_accuracy: 0.8477 - val_loss: 0.5082
Epoch 19/100
25/26 _____ 0s 22ms/step - accuracy: 0.9219 - loss:
0.2738
Epoch 19: val_accuracy did not improve from 0.84770
26/26 _____ 1s 33ms/step - accuracy: 0.9218 - loss:
0.2722 - val_accuracy: 0.8477 - val_loss: 0.5157
Epoch 20/100
26/26 _____ 0s 24ms/step - accuracy: 0.9297 - loss:
0.2062
Epoch 20: val_accuracy improved from 0.84770 to 0.85920, saving model
to best_model_t1.keras
26/26 _____ 2s 49ms/step - accuracy: 0.9297 - loss:
0.2063 - val_accuracy: 0.8592 - val_loss: 0.5220
Epoch 21/100
24/26 _____ 0s 24ms/step - accuracy: 0.9489 - loss:
0.1606
Epoch 21: val_accuracy did not improve from 0.85920
26/26 _____ 1s 34ms/step - accuracy: 0.9475 - loss:
0.1644 - val_accuracy: 0.8506 - val_loss: 0.5491
Epoch 22/100
25/26 _____ 0s 23ms/step - accuracy: 0.9496 - loss:
0.1773
Epoch 22: val_accuracy did not improve from 0.85920
26/26 _____ 1s 33ms/step - accuracy: 0.9495 - loss:
0.1768 - val_accuracy: 0.8534 - val_loss: 0.5619
Epoch 23/100
25/26 _____ 0s 23ms/step - accuracy: 0.9426 - loss:
0.1617
Epoch 23: val_accuracy did not improve from 0.85920
26/26 _____ 1s 33ms/step - accuracy: 0.9423 - loss:
0.1627 - val_accuracy: 0.8506 - val_loss: 0.5470
Epoch 24/100
25/26 _____ 0s 22ms/step - accuracy: 0.9480 - loss:
0.1690
Epoch 24: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9478 - loss:
0.1691 - val_accuracy: 0.8477 - val_loss: 0.5443
```

```

Epoch 25/100
25/26 _____ 0s 22ms/step - accuracy: 0.9381 - loss:
0.2156
Epoch 25: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9390 - loss:
0.2122 - val_accuracy: 0.8420 - val_loss: 0.5659
Epoch 26/100
25/26 _____ 0s 22ms/step - accuracy: 0.9519 - loss:
0.1638
Epoch 26: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9515 - loss:
0.1641 - val_accuracy: 0.8563 - val_loss: 0.5769
Epoch 27/100
25/26 _____ 0s 22ms/step - accuracy: 0.9451 - loss:
0.1543
Epoch 27: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9457 - loss:
0.1533 - val_accuracy: 0.8448 - val_loss: 0.5686
Epoch 28/100
25/26 _____ 0s 23ms/step - accuracy: 0.9573 - loss:
0.1228
Epoch 28: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9574 - loss:
0.1234 - val_accuracy: 0.8534 - val_loss: 0.5494
Epoch 29/100
25/26 _____ 0s 22ms/step - accuracy: 0.9319 - loss:
0.1751
Epoch 29: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9328 - loss:
0.1737 - val_accuracy: 0.8506 - val_loss: 0.5280
Epoch 30/100
25/26 _____ 0s 23ms/step - accuracy: 0.9540 - loss:
0.1456
Epoch 30: val_accuracy did not improve from 0.85920
26/26 _____ 1s 32ms/step - accuracy: 0.9538 - loss:
0.1462 - val_accuracy: 0.8592 - val_loss: 0.5285
Epoch 30: early stopping

```

##Print metrics

```

from tensorflow.keras.models import load_model
model_t1 = load_model('best_model_t1.keras')
print_scores(model_t1, val_ds, train_ds, 'Dropout Adjusted Transfer
Model')
# plot_accuracy_loss_graphs(history_t1, 'Dropout adjusted transfer
model')

```

Displaying accuracy and loss for Dropout Adjusted Transfer Model
Test loss:0.5220

```
Test accuracy:0.8592
Train loss:0.0325
train accuracy:0.9963
```

Increasing dropout to 0.4 has reduced overfitting a little

##Augmented Transfer Model

```
data_augmentation =
[ layers.RandomFlip("horizontal",input_shape=(height,width,3)),layers.R
andomRotation(0.1),layers.RandomZoom(0.1)]
model_augmented =
create_transfer_model(mob_model,dropout=0.4,data_augmentation=data_aug
mentation)
```

```
added <RandomFlip name=random_flip, built=False> to model
added <RandomRotation name=random_rotation, built=False> to model
added <RandomZoom name=random_zoom, built=False> to model
```

Model: "sequential_4"

Layer (type) Param #	Output Shape
random_flip (RandomFlip) 0	(None, 224, 224, 3)
random_rotation (RandomRotation) 0	(None, 224, 224, 3)
random_zoom (RandomZoom) 0	(None, 224, 224, 3)
rescaling_4 (Rescaling) 0	(None, 224, 224, 3)
MobileNet (Functional) 3,228,864	(None, 7, 7, 1024)
global_average_pooling2d_4 0	(None, 1024)
(GlobalAveragePooling2D)	

0	dropout_11 (Dropout)	(None, 1024)
0	flatten_4 (Flatten)	(None, 1024)
262,400	dense_12 (Dense)	(None, 256)
1,024	batch_normalization_4 (BatchNormalization)	(None, 256)
0	dropout_12 (Dropout)	(None, 256)
32,896	dense_13 (Dense)	(None, 128)
512	batch_normalization_5 (BatchNormalization)	(None, 128)
0	dropout_13 (Dropout)	(None, 128)
1,161	dense_14 (Dense)	(None, 9)

Total params: 3,526,857 (13.45 MB)

Trainable params: 297,225 (1.13 MB)

Non-trainable params: 3,229,632 (12.32 MB)

##Compile the model

```
model_augmented.compile(optimizer='adam',  
loss='sparse_categorical_crossentropy', metrics=['accuracy'])  
print('model_augmented compiled')
```

model_augmented compiled

##Train the model

```
early_stop=EarlyStopping(monitor='val_accuracy',mode='max',verbose=1,p  
atience=10,min_delta=0.01)  
checkpoint=ModelCheckpoint('best_model_a.keras',monitor='val_accuracy'  
,mode='max',verbose=1,save_best_only=True)
```

```
epochs=100  
history_a = model_augmented.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=epochs,  
    batch_size = batch_size,  
    callbacks = [early_stop,checkpoint]  
)
```

histories.append(history_a)

Epoch 1/100

26/26 _____ 0s 59ms/step - accuracy: 0.1796 - loss: 2.9263

Epoch 1: val_accuracy improved from -inf to 0.49425, saving model to best_model_a.keras

26/26 _____ 11s 143ms/step - accuracy: 0.1832 - loss: 2.9111 - val_accuracy: 0.4943 - val_loss: 1.5626

Epoch 2/100

25/26 _____ 0s 44ms/step - accuracy: 0.4971 - loss: 1.6661

Epoch 2: val_accuracy improved from 0.49425 to 0.67529, saving model to best_model_a.keras

26/26 _____ 2s 80ms/step - accuracy: 0.4999 - loss: 1.6569 - val_accuracy: 0.6753 - val_loss: 1.2297

Epoch 3/100

25/26 _____ 0s 43ms/step - accuracy: 0.6437 - loss: 1.1335

Epoch 3: val_accuracy improved from 0.67529 to 0.72126, saving model to best_model_a.keras

26/26 _____ 2s 80ms/step - accuracy: 0.6434 - loss: 1.1357 - val_accuracy: 0.7213 - val_loss: 0.9651

Epoch 4/100

25/26 _____ 0s 43ms/step - accuracy: 0.6995 - loss: 1.0396

Epoch 4: val_accuracy improved from 0.72126 to 0.75862, saving model to best_model_a.keras
26/26 _____ 2s 72ms/step - accuracy: 0.6987 - loss: 1.0411 - val_accuracy: 0.7586 - val_loss: 0.7938
Epoch 5/100
25/26 _____ 0s 44ms/step - accuracy: 0.7338 - loss: 0.8352
Epoch 5: val_accuracy improved from 0.75862 to 0.79023, saving model to best_model_a.keras
26/26 _____ 2s 73ms/step - accuracy: 0.7333 - loss: 0.8383 - val_accuracy: 0.7902 - val_loss: 0.7027
Epoch 6/100
26/26 _____ 0s 48ms/step - accuracy: 0.7402 - loss: 0.8620
Epoch 6: val_accuracy improved from 0.79023 to 0.80460, saving model to best_model_a.keras
26/26 _____ 3s 83ms/step - accuracy: 0.7401 - loss: 0.8614 - val_accuracy: 0.8046 - val_loss: 0.6599
Epoch 7/100
26/26 _____ 0s 45ms/step - accuracy: 0.7661 - loss: 0.8005
Epoch 7: val_accuracy did not improve from 0.80460
26/26 _____ 2s 63ms/step - accuracy: 0.7661 - loss: 0.8008 - val_accuracy: 0.7902 - val_loss: 0.6540
Epoch 8/100
25/26 _____ 0s 44ms/step - accuracy: 0.7380 - loss: 0.7764
Epoch 8: val_accuracy did not improve from 0.80460
26/26 _____ 2s 62ms/step - accuracy: 0.7401 - loss: 0.7736 - val_accuracy: 0.8046 - val_loss: 0.6129
Epoch 9/100
25/26 _____ 0s 44ms/step - accuracy: 0.7636 - loss: 0.8041
Epoch 9: val_accuracy did not improve from 0.80460
26/26 _____ 3s 61ms/step - accuracy: 0.7645 - loss: 0.7990 - val_accuracy: 0.8017 - val_loss: 0.6112
Epoch 10/100
25/26 _____ 0s 44ms/step - accuracy: 0.7714 - loss: 0.7370
Epoch 10: val_accuracy improved from 0.80460 to 0.83046, saving model to best_model_a.keras
26/26 _____ 2s 73ms/step - accuracy: 0.7727 - loss: 0.7311 - val_accuracy: 0.8305 - val_loss: 0.5695
Epoch 11/100
25/26 _____ 0s 44ms/step - accuracy: 0.8146 - loss: 0.5684
Epoch 11: val_accuracy improved from 0.83046 to 0.83908, saving model to best_model_a.keras
26/26 _____ 2s 73ms/step - accuracy: 0.8141 - loss:

```
0.5716 - val_accuracy: 0.8391 - val_loss: 0.5502
Epoch 12/100
26/26 _____ 0s 47ms/step - accuracy: 0.8190 - loss:
0.5820
Epoch 12: val_accuracy improved from 0.83908 to 0.84770, saving model
to best_model_a.keras
26/26 _____ 2s 88ms/step - accuracy: 0.8187 - loss:
0.5838 - val_accuracy: 0.8477 - val_loss: 0.5192
Epoch 13/100
26/26 _____ 0s 49ms/step - accuracy: 0.8158 - loss:
0.5652
Epoch 13: val_accuracy improved from 0.84770 to 0.85057, saving model
to best_model_a.keras
26/26 _____ 2s 78ms/step - accuracy: 0.8155 - loss:
0.5666 - val_accuracy: 0.8506 - val_loss: 0.5072
Epoch 14/100
25/26 _____ 0s 44ms/step - accuracy: 0.8519 - loss:
0.5318
Epoch 14: val_accuracy did not improve from 0.85057
26/26 _____ 2s 62ms/step - accuracy: 0.8511 - loss:
0.5323 - val_accuracy: 0.8333 - val_loss: 0.5093
Epoch 15/100
25/26 _____ 0s 45ms/step - accuracy: 0.8508 - loss:
0.4784
Epoch 15: val_accuracy did not improve from 0.85057
26/26 _____ 2s 63ms/step - accuracy: 0.8488 - loss:
0.4832 - val_accuracy: 0.8391 - val_loss: 0.5050
Epoch 16/100
25/26 _____ 0s 45ms/step - accuracy: 0.8357 - loss:
0.5292
Epoch 16: val_accuracy did not improve from 0.85057
26/26 _____ 2s 63ms/step - accuracy: 0.8354 - loss:
0.5294 - val_accuracy: 0.8448 - val_loss: 0.4983
Epoch 17/100
25/26 _____ 0s 45ms/step - accuracy: 0.8329 - loss:
0.4961
Epoch 17: val_accuracy did not improve from 0.85057
26/26 _____ 2s 63ms/step - accuracy: 0.8341 - loss:
0.4923 - val_accuracy: 0.8362 - val_loss: 0.4988
Epoch 18/100
25/26 _____ 0s 44ms/step - accuracy: 0.8718 - loss:
0.4101
Epoch 18: val_accuracy did not improve from 0.85057
26/26 _____ 2s 61ms/step - accuracy: 0.8706 - loss:
0.4141 - val_accuracy: 0.8448 - val_loss: 0.5060
Epoch 19/100
26/26 _____ 0s 47ms/step - accuracy: 0.8665 - loss:
0.3765
Epoch 19: val_accuracy did not improve from 0.85057
```

```

26/26 _____ 2s 65ms/step - accuracy: 0.8663 - loss:
0.3776 - val_accuracy: 0.8391 - val_loss: 0.5282
Epoch 20/100
26/26 _____ 0s 45ms/step - accuracy: 0.8510 - loss:
0.4585
Epoch 20: val_accuracy improved from 0.85057 to 0.85345, saving model
to best_model_a.keras
26/26 _____ 3s 75ms/step - accuracy: 0.8507 - loss:
0.4588 - val_accuracy: 0.8534 - val_loss: 0.5163
Epoch 21/100
25/26 _____ 0s 44ms/step - accuracy: 0.8501 - loss:
0.4077
Epoch 21: val_accuracy did not improve from 0.85345
26/26 _____ 2s 62ms/step - accuracy: 0.8486 - loss:
0.4112 - val_accuracy: 0.8506 - val_loss: 0.4837
Epoch 22/100
25/26 _____ 0s 44ms/step - accuracy: 0.8531 - loss:
0.4400
Epoch 22: val_accuracy did not improve from 0.85345
26/26 _____ 3s 62ms/step - accuracy: 0.8532 - loss:
0.4389 - val_accuracy: 0.8506 - val_loss: 0.4906
Epoch 22: early stopping

```

##Print accuracy and loss curves

```

from tensorflow.keras.models import load_model
model_a = load_model('best_model_a.keras')
print_scores(model_t1, val_ds, train_ds, 'Augmented Transfer Model')
# plot_accuracy_loss_graphs(history_a, 'Augmented transfer model')

Displaying accuracy and loss for Augmented Transfer Model
Test loss:0.5220
Test accuracy:0.8592
Train loss:0.0325
train accuracy:0.9963

```

##Plot loss/accuracy curves for all models

The final accuracy and loss are the same for augmented ad non augmented models. The augmented model has a slightly lesser ovetfitting

```


import matplotlib.image as mpimg
histories = [history_c, history_t, history_t1, history_a]
#titles = ['CNN Model', 'Transfer Learning model using
MobileNet', 'Dropout adjusted transfer model', 'Augmented transfer
model']
#print(test_accuracy.keys())

titles = list(test_accuracy.keys())


```

```
for i,history in enumerate(histories):
    plot_accuracy_loss_graphs(history,titles[i],show=False)
plt.figure(figsize=(16,16))
for i,title in enumerate(titles):
    title = titles[i]
    plt.subplot(2,2,i+1)
    figure = title+'.png'
    img = mpimg.imread(figure)
    plt.imshow(img)
    plt.axis('off')
    accuracy = test_accuracy[title]
    loss = test_loss[title]
    title = title + f'\n(validation accuracy = {accuracy:.4f}, loss =
{loss:.4f})'
    plt.title(title)
plt.show()
```

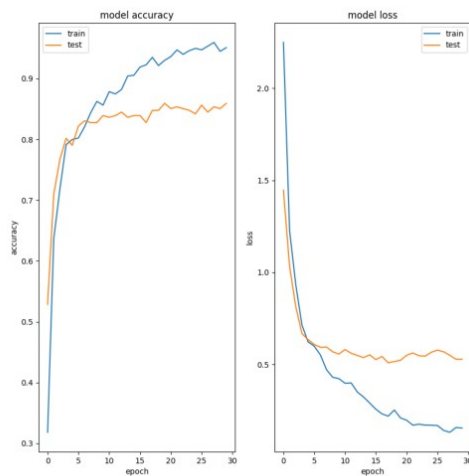
CNN Model
(validation accuracy = 0.6121, loss = 1.2901)



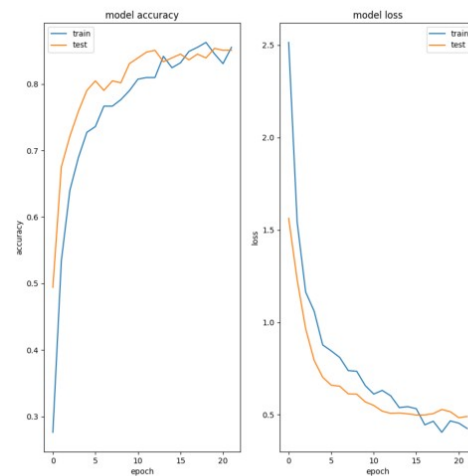
Transfer Model
(validation accuracy = 0.8621, loss = 0.5267)



Dropout Adjusted Transfer Model
(validation accuracy = 0.8592, loss = 0.5220)



Augmented Transfer Model
(validation accuracy = 0.8592, loss = 0.5220)




##Visualize the predictions

```
image_map = get_predicted_labels(model_a,class_names)
print('size of image map = ',len(image_map))
display_test_images(image_map,'Transfer Learning model using
MobileNet. Title is the prediction')
```

size of image map = 300

Predictions for Transfer Learning model using MobileNet. Title is the prediction



The predictions with the transfer model is much more accurate (expected accuracy is 86%) when compared to the CNN model. We can see that even classes with very few train images have been correctly identified(paper boat for example)

##Metrics

```
true_labels, predicted_labels =
get_metrics(model_a, val_ds, val_split=0.3)
```

##Confusion Matrix


```

from sklearn.metrics import ConfusionMatrixDisplay
cm = confusion_matrix(true_labels,predicted_labels)
print('Confusion Matrix')
print(cm)
print()

print('Confusion Matrix Displayed\n')
plt.figure(figsize = (10,10))
disp =
ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=class_names)
disp.plot()
plt.title('Confusion Matrixr')
plt.xticks(rotation=90)
plt.show()

```

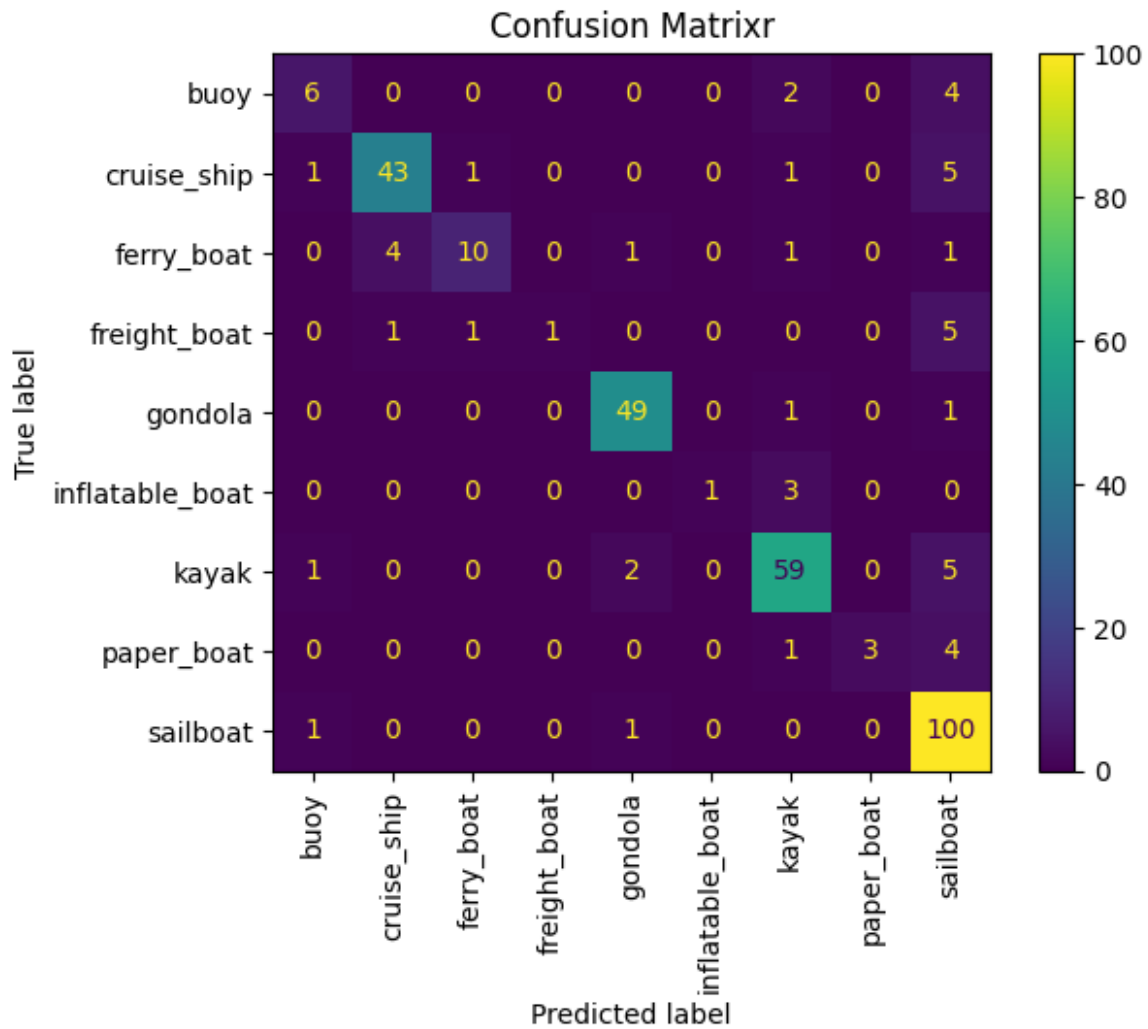
```

Confusion Matrix
[[ 6  0  0  0  0  0  2  0  4]
 [ 1 43  1  0  0  0  1  0  5]
 [ 0  4 10  0  1  0  1  0  1]
 [ 0  1  1  1  0  0  0  0  5]
 [ 0  0  0  0 49  0  1  0  1]
 [ 0  0  0  0  0  1  3  0  0]
 [ 1  0  0  0  2  0 59  0  5]
 [ 0  0  0  0  0  0  1  3  4]
 [ 1  0  0  0  1  0  0  0 100]]

```

Confusion Matrix Displayed

<Figure size 1000x1000 with 0 Axes>



##Classification Report

```
from sklearn.metrics import classification_report

class_name_dict = {
    0 : ''
}

cl_report = classification_report(true_labels,predicted_labels)
print(cl_report)

print('Index to name mapping:\n')

for k,v in class_dict.items():
    print(f'{k} - > {v}')
```

	precision	recall	f1-score	support
0	0.67	0.50	0.57	12

1	0.90	0.84	0.87	51
2	0.83	0.59	0.69	17
3	1.00	0.12	0.22	8
4	0.92	0.96	0.94	51
5	1.00	0.25	0.40	4
6	0.87	0.88	0.87	67
7	1.00	0.38	0.55	8
8	0.80	0.98	0.88	102
accuracy			0.85	320
macro avg	0.89	0.61	0.67	320
weighted avg	0.86	0.85	0.83	320
Index to name mapping:				
0 - > inflatable_boat				
1 - > freight_boat				
2 - > paper_boat				
3 - > buoy				
4 - > ferry_boat				
5 - > cruise_ship				
6 - > gondola				
7 - > kayak				
8 - > sailboat				

As expected the classes that have higher support(more train images) have better overall scores

Inferences

##CNN Model

The CNN model built in part 1 of this project gave an accuracy of around 65%. When tested on the 300 test images, we could see many mis-classifications. This could be because the training set contained just 1162 images and the code used only 80% of that for training. This is a very low number for the model to learn well. The model has a bit of overfitting, though :

##Transfer Learning

Accuracy significantly increases with transfer model (Mobilenet_V2) with fine tuning with custom FCN. However there is significant over fitting with low dropout values. With higher dropout values, overfitting gets reduced without sacrificing accuracy Higher dropout value with additional image augmentation reduces overfitting further at every epoch without compromising accuracy The higher accuracy is evident in the number of correct predictions. This can also be seen in the confusion matrix and classification report. This has been verified visually too on the unlabelled test data set

##Conclusion

For custom datasets that do not have sufficiently large number of images, transfer learning is the preferred method of choice. By freezing the model's weights and then fine tuning with a fully connected layer gives us a very accurate model that otherwise would not have been possible.